

# 5042259201

## Big Data Analytics

1 Week Part 1

Introduction to Big Data Mining  
& Hadoop



# Introduction to Class (1/3)

---

## Big Data Analytics

### ■ Lecture Web site ; Smart Campus

- Soongsil campus lecture web site

### ■ Lecture Objectives

- To study Big data Analytics algorithms, tools, and programming.
- To discuss application of Big data Analytics algorithms to real world problems



# Introduction to Class(2/3)

## ■ Textbook:



- “Mining of Massive Datasets,” J. Leskovec, et al.  
<http://www.mmnds.org>

## ■ Auxiliary Textbooks

- “Advanced Analytics with Spark,” S. Ryza, et al., O’Reilly
- “Hadoop: The Definitive Guide,” Tom White, 4<sup>th</sup> ed., O’Reilly
- “Spark: The Definitive Guide, “ B. chambers, et. al., O’Reilly



Sandy Ryza, Uri Laserson,  
Sean Owen, & Josh Wills



Bill Chambers & Matei Zaharia

# Introduction to Class(3/3)

---

- Class Style

- Students' Presentations (+ Lecture)

- Evaluation

- Presentation(50%) – Present Chosen Subjects  
from Textbooks

- Final Project(50%)



# Tentative Lecture Schedule

Week	Keyword	Description
01	Data Mining, MapReduce, Spark Programming (Scala, Pyspark), Tensorflow	Introduction to Data Mining, Statistical Modeling, Machine Learning, MapReduce algorithms, Spark, PySpark
02	Spark Toolsets, Spark Structured API Spark Machine Learning	RDD, DataFrame, Dataset, Spark MLlib, Spark ML
03	Similarity	Finding Similar Items, Locality-Sensitive Functions, Hashing
04	Mining Data Streams	Stream Data Model, Sampling Data in a Stream, Filtering Streams
05	Link Analysis	PageRank, Link Spam, Hubs and Authorities
06	Frequent Itemsets	Market-Basket Model, A-Priori Algorithm, Limited-Pass Algorithms
07	Clustering	Hierarchical Clustering, K-Means Clustering, CUR Algorithm
08	Advertising on the Web	Issues in On-Line Advertising, Matching Problem, Adwords Problem
09	Recommendation System	Collaborative Filtering, Dimensionality Reduction, Netflix Challenge
10	Mining Social-Network Graphs	Mining Social-Network, Graphs
11	Dimensionality Reduction	PCA, SVD, CUR Decomposition
12	Large-Scale Machine Learning	SVM. Decision Tree
13	Project Presentations	Project Presentations
14	Project Presentations	Project Presentations
15	Project Presentations	Project Presentations

# Contents of 1<sup>st</sup> Week

---

- What is Data Mining?
- Big Data Computing Framework
- Distributed File System
- MapReduce
- Introduction to Apache Hadoop
- Hadoop: HDFS
- Hadoop: MapReduce
- Apache YARN
- Introduction to Apache Spark



# **What is Data Mining?**

## **Knowledge discovery from data**



Chapter 1,  
“Mining of Massive Datasets, “  
Jure Leskovec, Anand Rajaraman,  
Jeff Ullman Stanford University  
<http://www.mmds.org>

# What is Data Mining?

---

- Given lots of data
- Discover patterns and models that are:
  - **Valid:** hold on new data with some certainty
  - **Useful:** should be possible to act on the item
  - **Unexpected:** non-obvious to the system
  - **Understandable:** humans should be able to interpret the pattern



# Data Mining

---

- But to extract the knowledge data needs to be
  - Stored
  - Managed
  - and Analyzed ← this class

Data Mining  $\approx$  Big Data  $\approx$   
Predictive Analytics  $\approx$  Data Science

# Data Mining Tasks

## -Business Analytics

### ■ Descriptive Analytics

- Looks at past performance and understands that performance by mining historical data to look for the reasons behind past success or failure.

### ■ Predictive Analytics

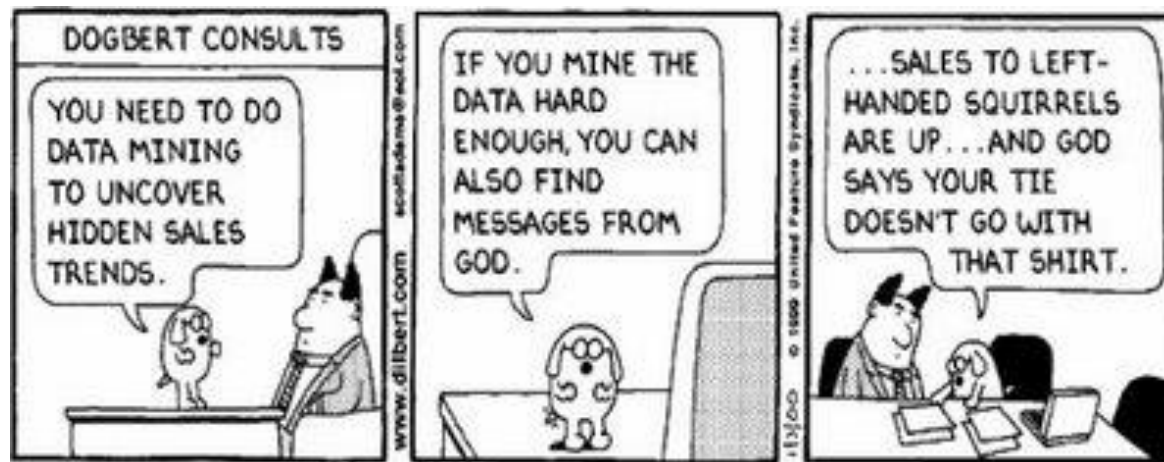
- Historical data is combined with rules, algorithms, and occasionally external data to determine the probable future outcome of an event or the likelihood of a situation occurring.

### ■ Prescriptive Analytics

- suggests decision options on how to take advantage of a future opportunity or mitigate a future risk and shows the implication of each decision option.

# Meaningfulness of Analytic Answers

- A risk with “Data mining” is that an analyst can “discover” patterns that are meaningless
- Statisticians call it **Bonferroni's principle**:
  - Roughly, if you look in more places for interesting patterns than your amount of data will support, you are bound to find crap



# Bonferroni's Principle

- Informal presentation of a statistical theorem
  - If your method of finding significant items returns significantly more items that you would expect in the actual population, you can assume most of the items you find with it are bogus.
  - This essentially means that an algorithm or method we think is useful for finding a particular set of data actually returns more false positives as it returns larger portions of the data than should be within that category.

# Ex. of Bonferroni's principle

- Assume you are trying to identify people who are cheating on their spouses within a certain population, and you know that the percentage in the population who cheat on their spouses is 5%.
- If you decide that people who claim to go out with coworkers more than three times a month are most likely actually cheating on their spouses, but discover that 20% of people in the population qualify with your method, then you know in the very best case only one quarter of the people you identify will actually be cheaters.
- Furthermore, if there are any false negatives (cheaters who aren't identified as cheaters), an even higher percentage of the "cheaters" identified with the system would be false positives.



# Uses

---

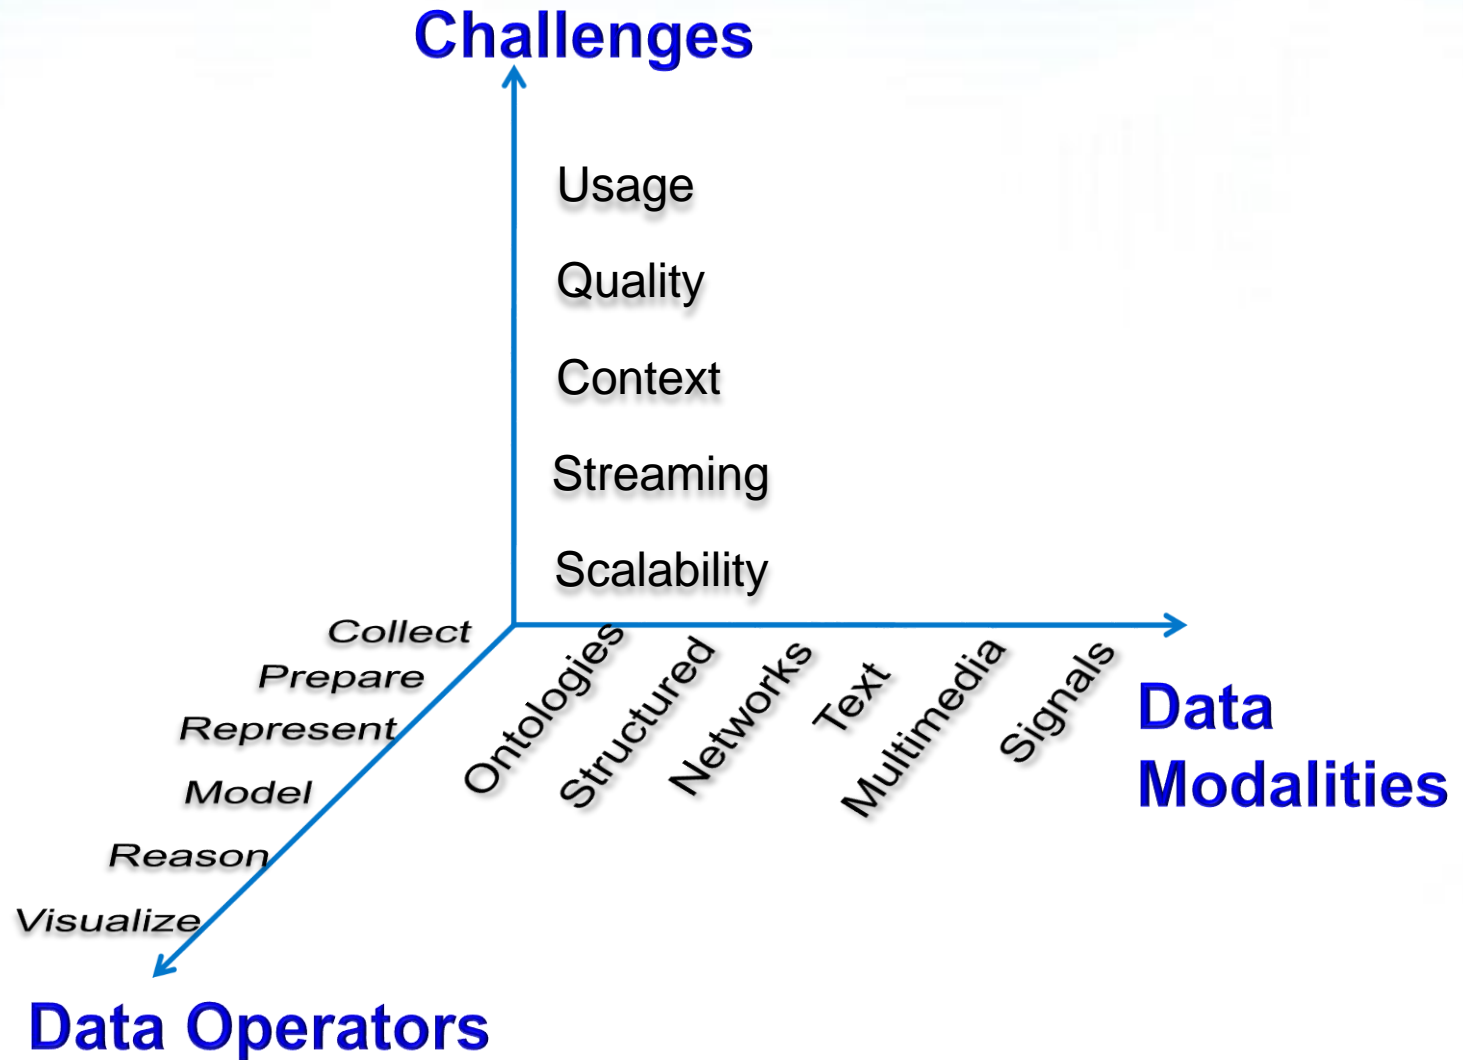
- Applying Bonferroni's Principle to an algorithm or system for identifying or classifying data gives you an upper bound on the accuracy of your methods. If you determine that you match significantly more data or less data than you should expect, then you in the best case have too many false positives or false negatives, respectively.
- The principle is especially useful in debunking individuals who use cold reading techniques. You may think they are using some sort of psychic power to accurately identify a single person, but if it turns out that 90% of the audience can identify with something they are saying and it's likely there's at least one person who can identify with most of what they are saying in every audience, their powers become much less impressive.

# Meaningfulness of Analytic Answers

Example:

- We want to find (unrelated) people who **at least twice have stayed at the same hotel on the same day**
  - $10^9$  people being tracked
  - 1,000 days
  - Each person stays in a hotel 1% of time (1 day out of 100)
  - Hotels hold 100 people (so  $10^5$  hotels)
  - If everyone behaves randomly (i.e., no terrorists) will the data mining detect anything suspicious?
- Expected number of “suspicious” pairs of people:
  - 250,000
  - ... too many combinations to check – we need to have some additional evidence to find “suspicious” pairs of people in some more efficient way

# What matters when dealing with data?



# Data Mining: Cultures


## ■ Data mining overlaps with:

- **Databases:** Large-scale data, simple queries
- **Machine learning:** Small data, Complex models
- **CS Theory:** (Randomized) Algorithms

## ■ Different cultures:


- To a DB person, data mining is an extreme form of **analytic processing** – queries that examine large amounts of data
  - Result is the query answer
- To a ML person, data-mining is the **inference of models**
  - Result is the parameters of the model

## ■ In this class we will do both!



# This Class

- This class overlaps with machine learning, statistics, artificial intelligence, databases but more stress on
  - Scalability (big data)
  - Algorithms
  - Computing architectures
  - Automation for handling large data





# What will we learn?

---

- We will learn to mine different types of data:
  - Data is high dimensional
  - Data is a graph
  - Data is infinite/never-ending
  - Data is labeled
- We will learn to use different models of computation:
  - MapReduce
  - Streams and online algorithms
  - Single machine in-memory

# What will we learn?

---

## ■ We will learn to solve real-world problems:

- Recommender systems
- Market Basket Analysis
- Spam detection
- Duplicate document detection

## ■ We will learn various “tools”:

- Linear algebra (SVD, Rec. Sys., Communities)
- Optimization (stochastic gradient descent)
- Dynamic programming (frequent itemsets)
- Hashing (LSH, Bloom filters)

# How It All Fits Together

High dim.  
data

Locality  
sensitive  
hashing

Clustering

Dimensional  
ity  
reduction

Graph  
data

PageRank,  
SimRank

Community  
Detection

Spam  
Detection

Infinite  
data

Filtering  
data  
streams

Web  
advertising

Queries on  
streams

Machine  
learning

SVM

Decision  
Trees

Perceptron,  
kNN

Apps

Recommen  
der systems

Association  
Rules

Duplicate  
document  
detection



**How do you want that data?**




# Big Data Computing Framework

Section 2.1,  
“Mining of Massive Datasets, “  
Jure Leskovec, Anand Rajaraman,  
Jeff Ullman Stanford University  
<http://www.mmds.org>



# Single Node Architecture

---




**Machine Learning, Statistics**

**“Classical” Data Mining**

# Motivation: Google Example

- 20+ billion web pages x 20KB = 400+ TB
- 1 computer reads 30–35 MB/sec from disk
  - ~4 months to read the web
- ~1,000 hard drives to store the web
- Takes even more to do something useful with the data!
- Today, a standard architecture for such problems is emerging:
  - Cluster of commodity Linux nodes
  - Commodity network (ethernet) to connect them

# Cluster Architecture



Each rack contains 16-64 nodes

In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>

# Large-scale Computing

- Large-scale computing for data mining problems on commodity hardware
- Challenges:
  - How do you distribute computation?
  - How can we make it easy to write distributed programs?
  - Machines fail:
    - One server may stay up 3 years (1,000 days)
    - If you have 1,000 servers, expect to loose 1/day
    - People estimated Google had ~1M machines in 2011
      - 1,000 machines fail every day!

# Idea and Solution

■ **Issue:** Copying data over a network takes time

■ **Idea:**

- Bring computation close to the data
- Store files multiple times for reliability

■ **MapReduce addresses these problems**

- Google's computational/data manipulation model
- Elegant way to work with big data
- **Storage Infrastructure – File system**
  - Google: GFS. Hadoop: HDFS
- **Programming model**
  - MapReduce





# Distributed File System

Section 2.1,  
“Mining of Massive Datasets,”  
Jure Leskovec, Anand Rajaraman,  
Jeff Ullman Stanford University  
<http://www.mmnds.org>

# Storage Infrastructure

---

## ■ Problem:

- If nodes fail, how to store data persistently?

## ■ Answer:

### ➤ Distributed File System:

- Provides global file namespace
- Google GFS; Hadoop HDFS;

## ■ Typical usage pattern

- Huge files (100s of GB to TB)
- Data is rarely updated in place
- Reads and appends are common

# Distributed File System

## ■ Chunk servers

- File is split into contiguous chunks
- Typically each chunk is 16–64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

## ■ Master node

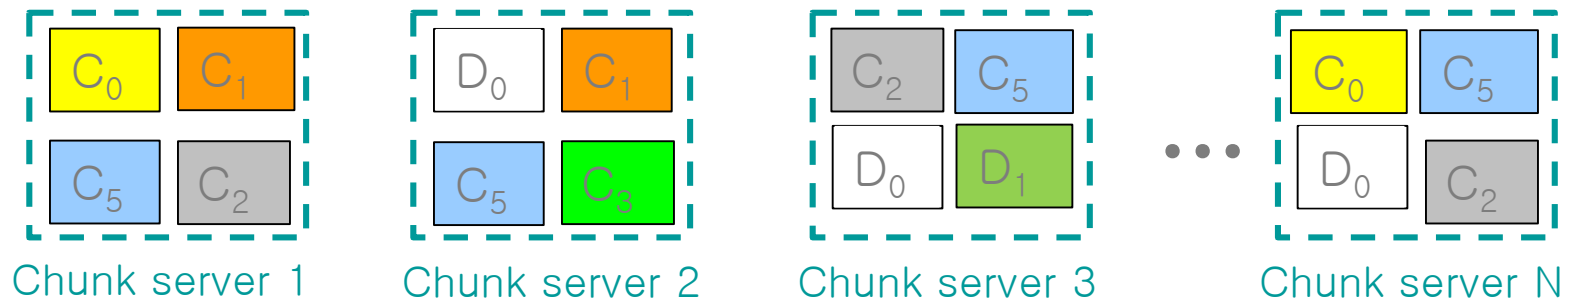
- a.k.a. Name Node in Hadoop's HDFS
- Stores metadata about where files are stored
- Might be replicated

## ■ Client library for file access

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

# Distributed File System

- Reliable distributed file system
- Data kept in “chunks” spread across machines
- Each chunk replicated on different machines
  - Seamless recovery from disk or machine failure



Bring computation directly to the data!

Chunk servers also serve as compute servers

# Distributed File System

---

- Don't move data to workers... move workers to the data!
  - Store data on the local disks of nodes in the cluster
  - Start up the workers on the node that has the data local
- Why?
  - Not enough RAM to hold all the data in memory
  - Disk access is slow, but disk throughput is reasonable
- A distributed file system is the answer
  - GFS (Google File System) for Google's MapReduce
  - HDFS (Hadoop Distributed File System) for Hadoop

# The Need for GFS

---

- Component failures normal
  - Due to clustered computing
- Files are huge
  - By traditional standards (many TB)
- Most mutations are mutations
  - Not random access overwrite
- Co-Designing apps & file system
- Typical: 1000 nodes & 300 TB



# Things to be desired

---

- Must monitor & recover from comp failures
- Modest number of large files
- Workload
  - Large streaming reads + small random reads
  - Many large sequential writes
    - Random access overwrites don't need to be efficient
- Need semantics for concurrent appends
- High sustained bandwidth
  - More important than low latency

# GFS: Assumptions

---

- Commodity hardware over “exotic” hardware
  - Scale “out”, not “up”
- High component failure rates
  - Inexpensive commodity components fail all the time
- “Modest” number of huge files
  - Multi-gigabyte files are common, if not encouraged
- Files are write-once, mostly appended to
  - Perhaps concurrently
- Large streaming reads over random access
  - High sustained throughput over low latency

# GFS: Design Decisions

---

- Files stored as chunks
  - Fixed size (64MB)
- Reliability through replication
  - Each chunk replicated across 3+ chunkservers
- Single master to coordinate access, keep metadata
  - Simple centralized management
- No data caching
  - Little benefit due to large datasets, streaming reads
- Simplify the API
  - Push some of the issues onto the client (e.g., data layout)

**HDFS = GFS clone (same basic ideas)**

# Interface

---

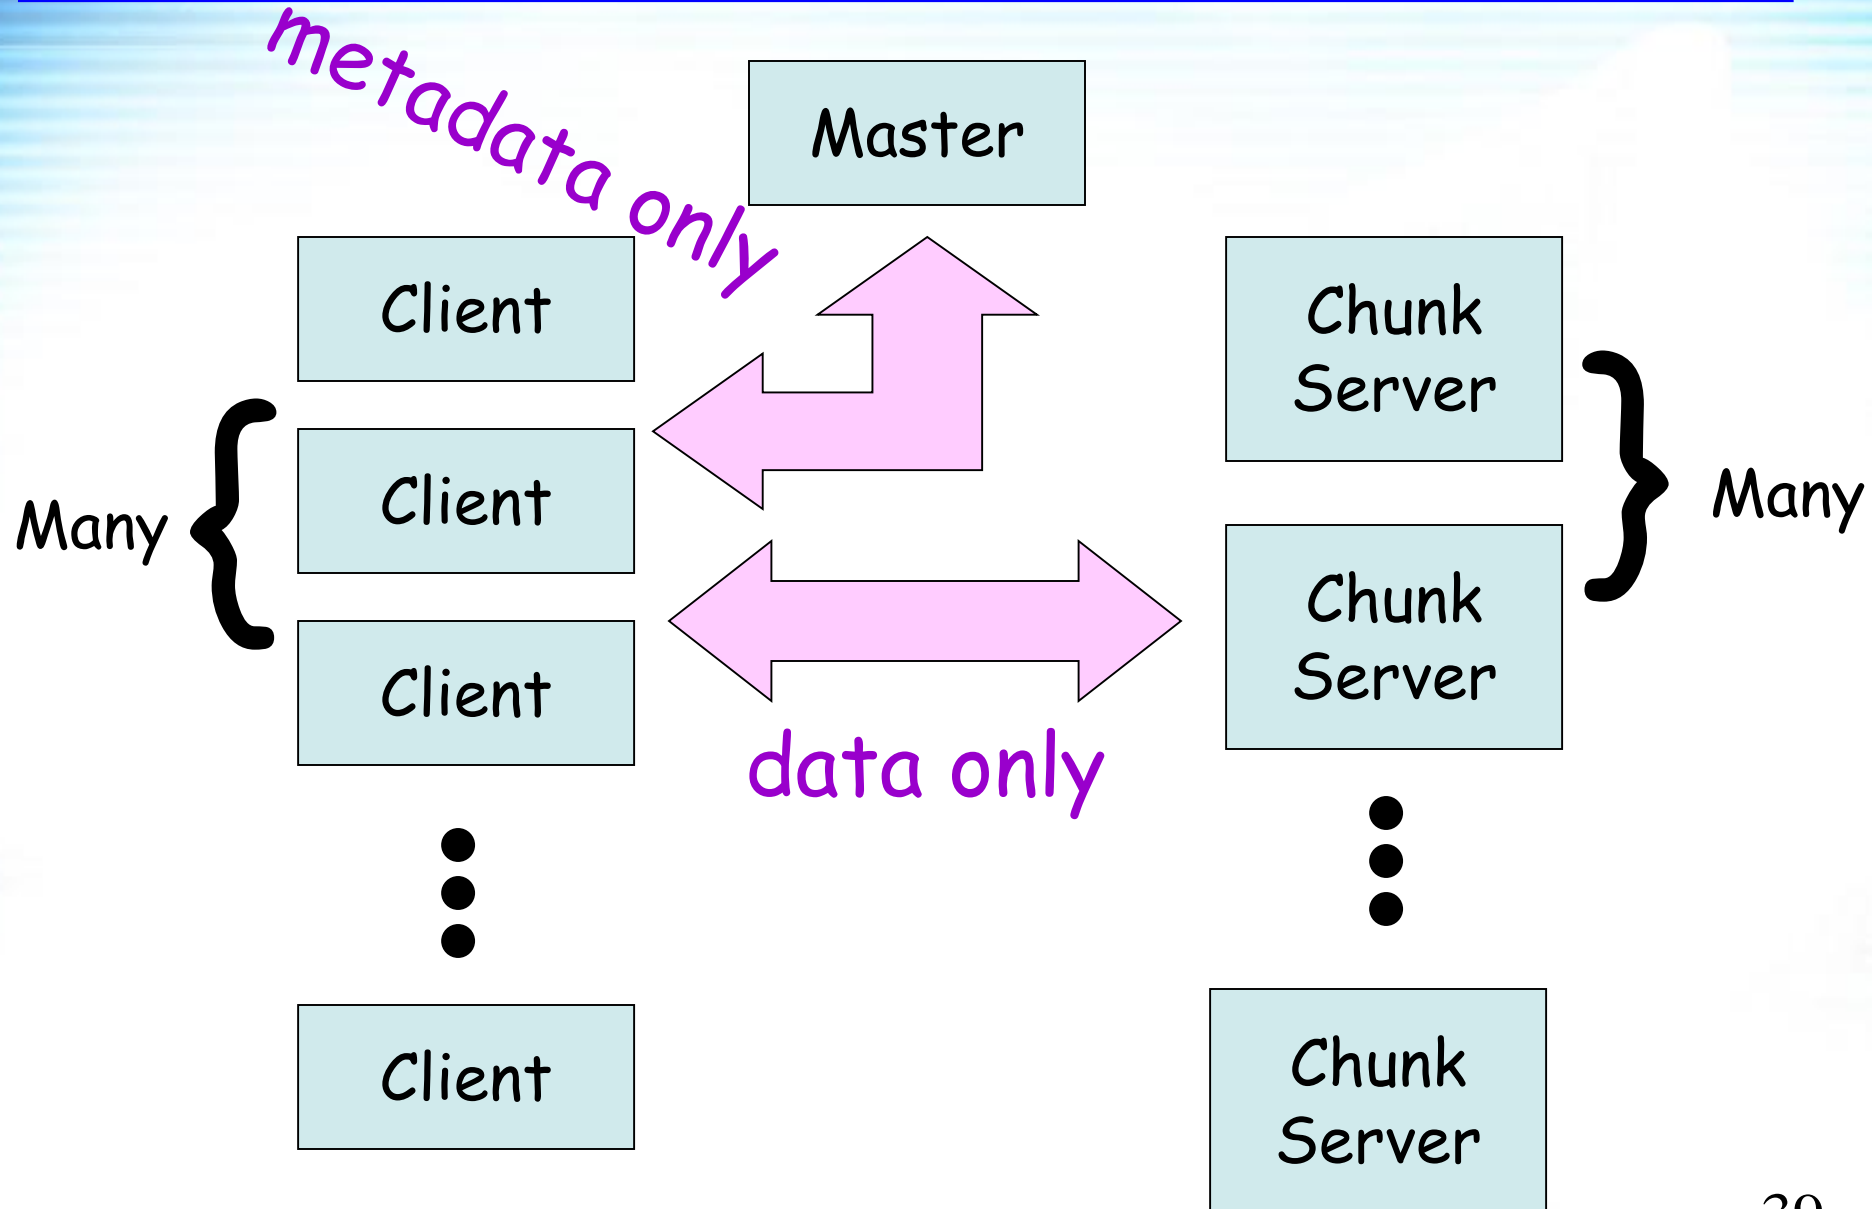
## ■ Familiar ; CRUD

- Create, Read, Update(write), Delete, Close

## ■ Novel

- Snapshot
  - Low cost
- Record append
  - Atomicity with multiple concurrent writes

# Architecture




# Architecture

## ■ Store all files

- In fixed-size chunks
  - 64 MB
  - 64 bit unique handle

## ■ Triple redundancy





# Architecture


---

Master

- Stores all metadata
  - Namespace
  - Access-control information
  - Chunk locations
  - 'Lease' management
- Heartbeats
- Having one master → global knowledge
  - Allows better placement / replication
  - Simplifies design

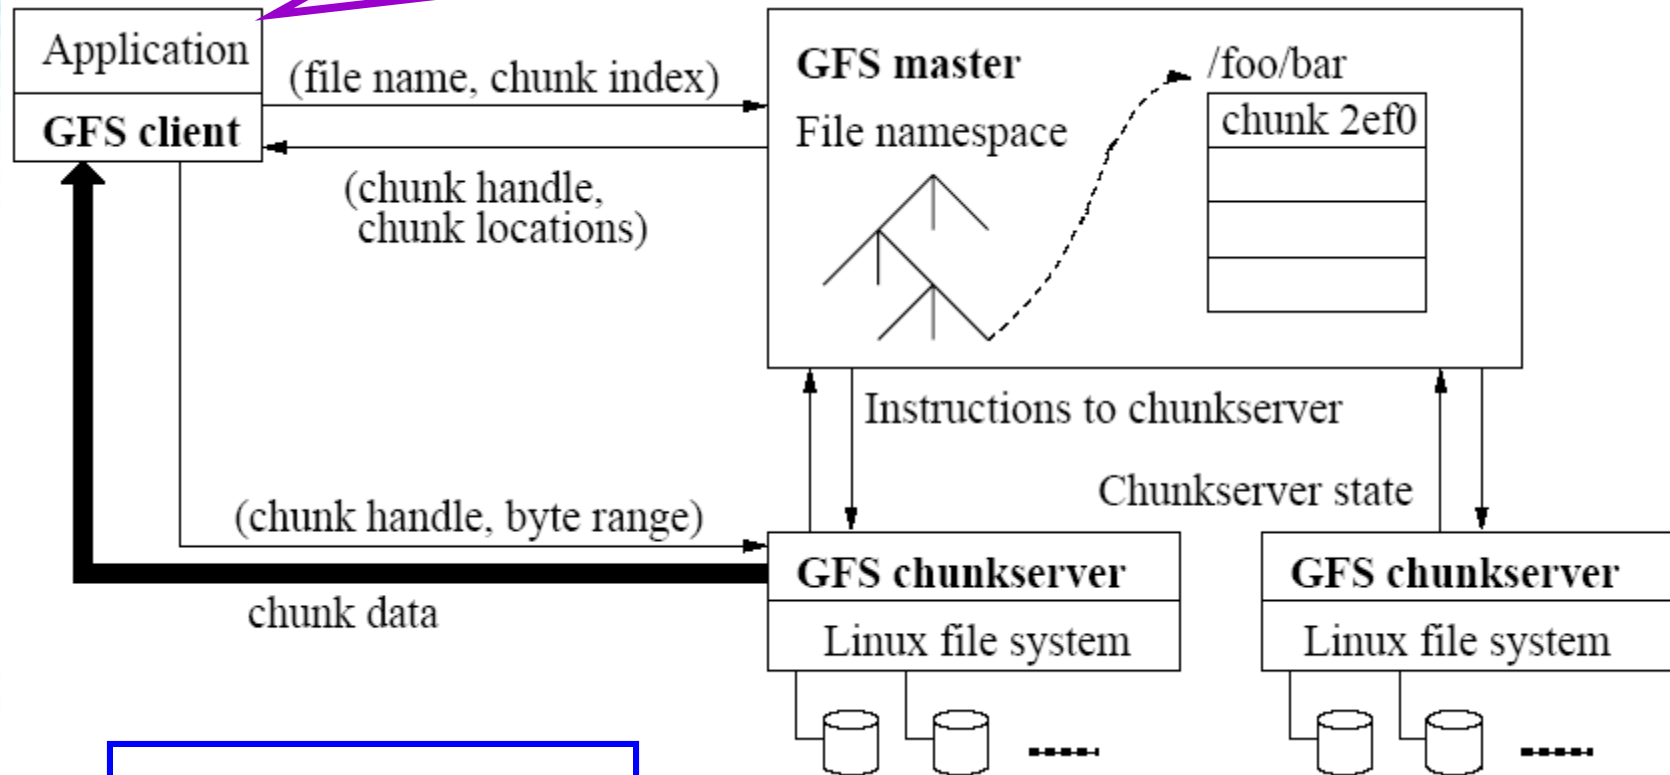
# Architecture

---



- GFS code implements API
- Cache only metadata

Using fixed chunk size, translate filename & byte offset to chunk index.  
Send request to master



Legend:

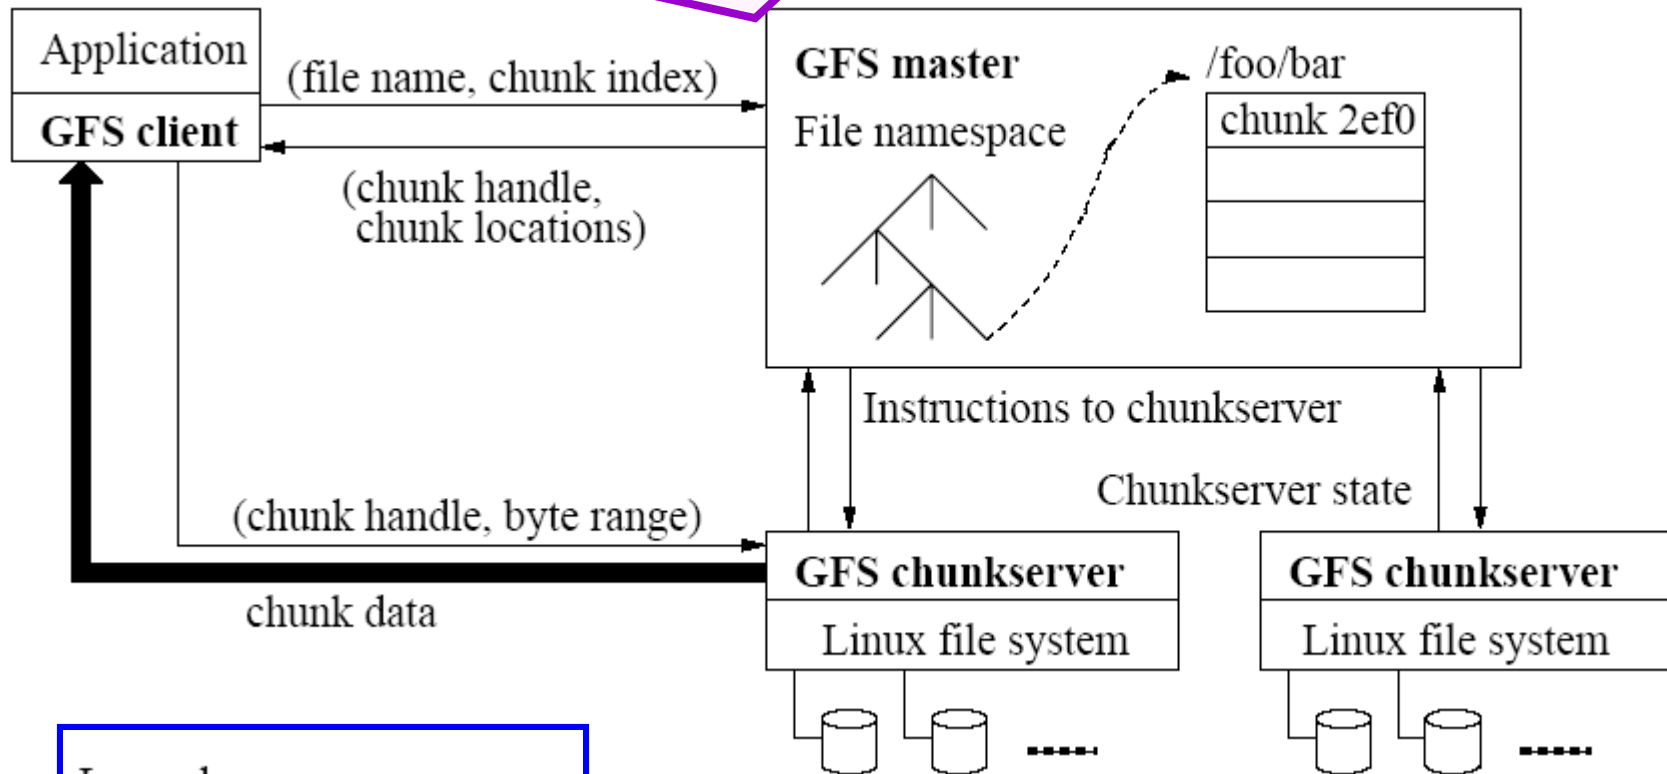


Data messages



Control messages

Replies with chunk handle & location of chunkserver replicas (including which is 'primary')



Legend:

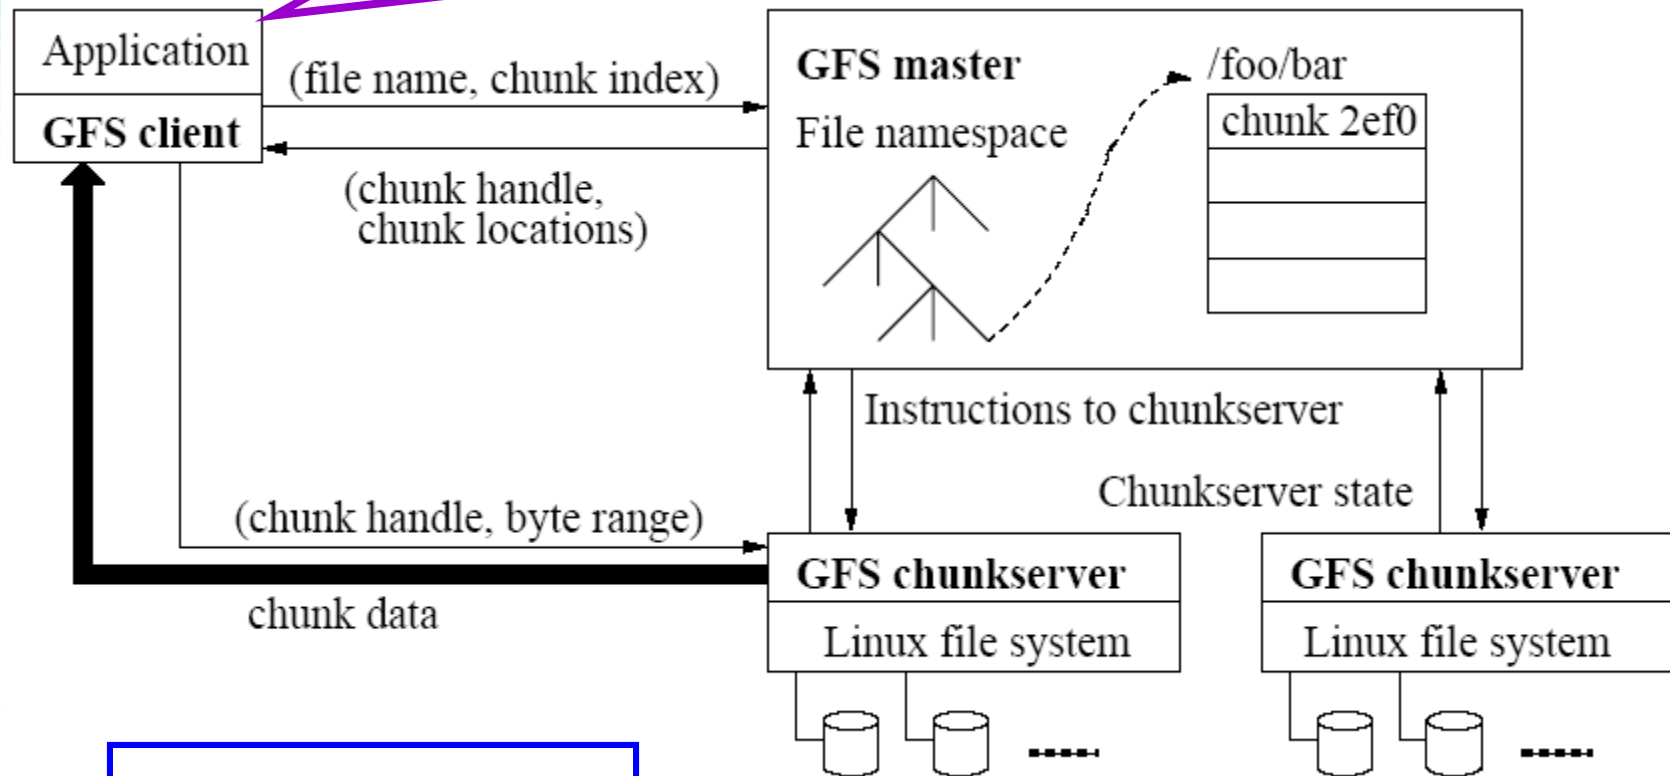


Data messages



Control messages

# Cache info using filename & chunk index as key



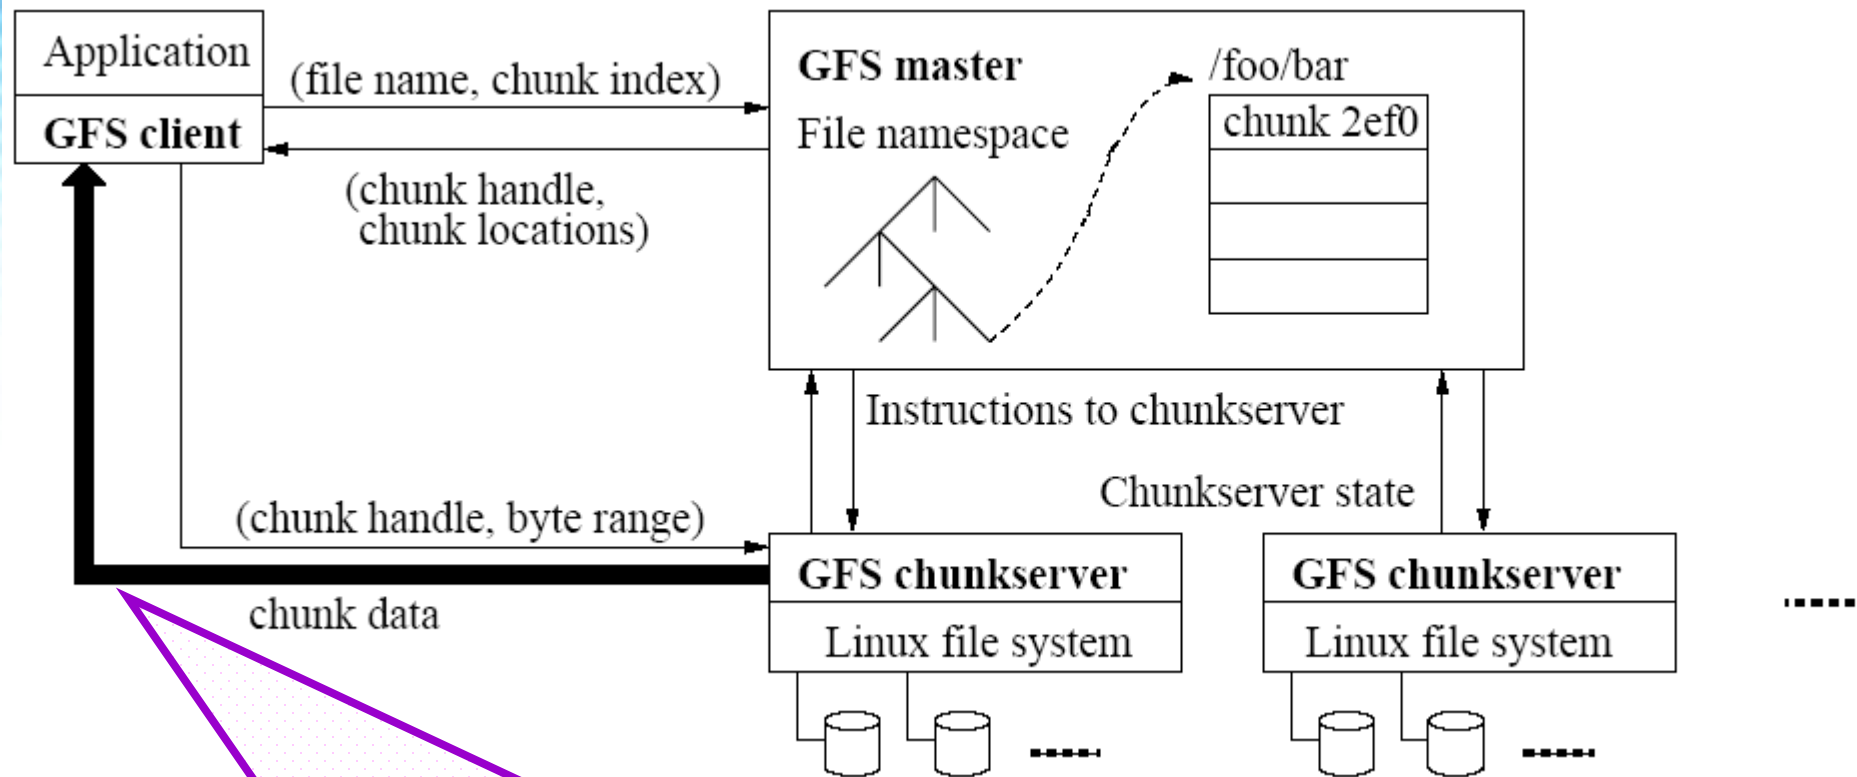
Legend:



Data messages



Control messages



Request data from nearest chunkserver  
"chunkhandle & index into chunk"



# Metadata

---

- Master stores three types
  - File & chunk namespaces
  - Mapping from files → chunks
  - Location of chunk replicas
- Stored in memory
- Kept persistent thru logging

# From GFS to HDFS

---

## ■ Terminology differences:

- GFS master = Hadoop namenode
- GFS chunkservers = Hadoop datanodes

## ■ Functional differences:

- HDFS performance is (likely) slower

For the most part, we'll use the Hadoop terminology...



# MapReduce

Section 2.2,  
“Mining of Massive Datasets, “  
Jure Leskovec, Anand Rajaraman,  
Jeff Ullman Stanford University  
<http://www.mmds.org>

# MapReduce

---

- Much of the course will be devoted to large scale computing for data mining
- Challenges:
  - How to distribute computation?
  - Distributed/parallel programming is hard
- **MapReduce** addresses all of the above
  - Google's computational/data manipulation model
  - Elegant way to work with big data

# Typical Large-Data Problem

*Map* ■ Iterate over a large number of records

■ Extract something of interest from each

■ Shuffle and sort intermediate results

■ Aggregate intermediate results *Reduce*


■ Generate final output

Key idea: provide a functional abstraction for these two operations

# Roots in Functional Programming

Map

Fold



**Functional programming +  
distributed computing!**

```
items = [1, 2, 3, 4, 5]
squared = list(map(lambda x: x**2, items))
from functools import reduce
summation=reduce((lambda x, y: x+y), squared)
print(summation)
```



# Functional Programming

- A programming paradigm, *i.e.* a style of computer programming that treats computation as the evaluation of mathematical functions
- Emphasizes functions that produce results that depend only on their inputs and not on the program state
  - *i.e.* pure mathematical functions
- In functional code, the output value of a function depends only on the arguments that are input to the function, so calling a function  $f$  twice with the same value for an argument  $x$  will produce the same result  $f(x)$  both times

# MapReduce

---

- A programming model for processing large data sets with a parallel, distributed algorithm on a cluster with massively parallel architecture
  - **Map()** procedure that performs functions such as filtering or sorting to each element of a container(e.g. a list), returning a container of results in the same order.
  - **Reduce()** procedure that performs a summary operation
    - *e.g.* counting the number of students in each queue, yielding surname frequencies

# MapReduce

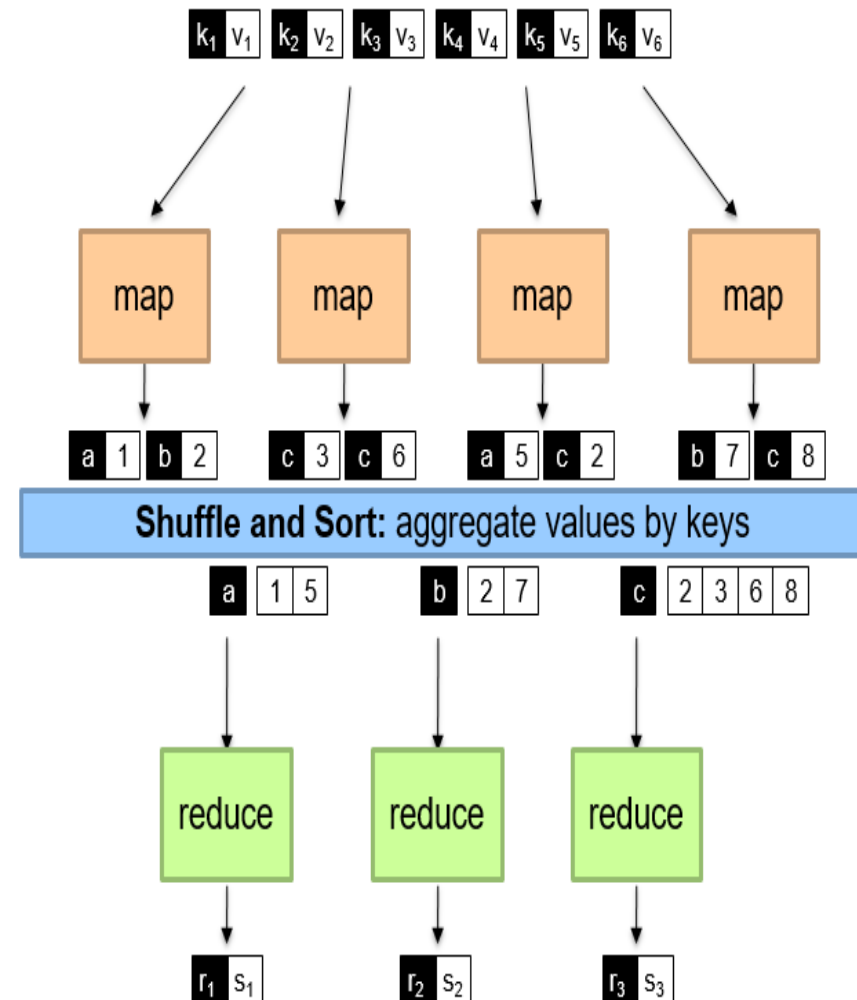
- Programmers specify two functions:

**map**  $(k_1, v_1) \rightarrow [ \langle k_2, v_2 \rangle ]$

**reduce**  $(k_2, [v_2]) \rightarrow [ \langle k_3, v_3 \rangle ]$

- All values with the same key are sent to the same reducer

- The execution framework handles everything else...



# Task: Word Count

---

## Case 1:

- File too large for memory, but all <word, count> pairs fit in memory

## Case 2:

### ■ Count occurrences of words:

- `words(doc.txt) | sort | uniq -c`

- where **words** takes a file and outputs the words in it, one per a line

### ■ Case 2 captures the essence of MapReduce

- Great thing is that it is naturally parallelizable

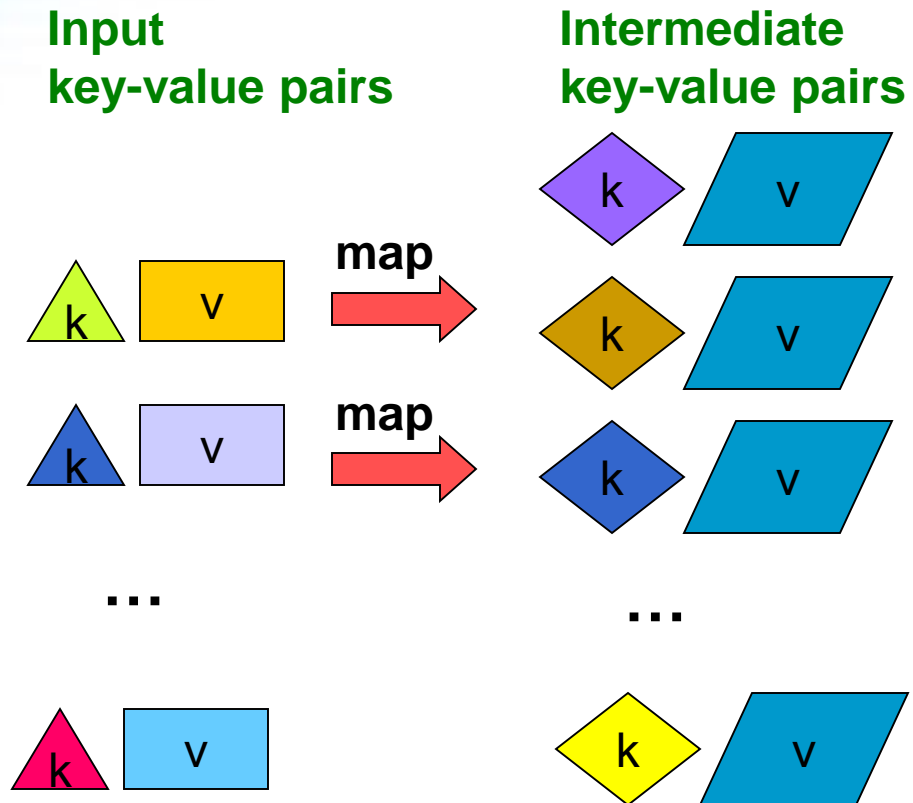
# MapReduce: Overview

---

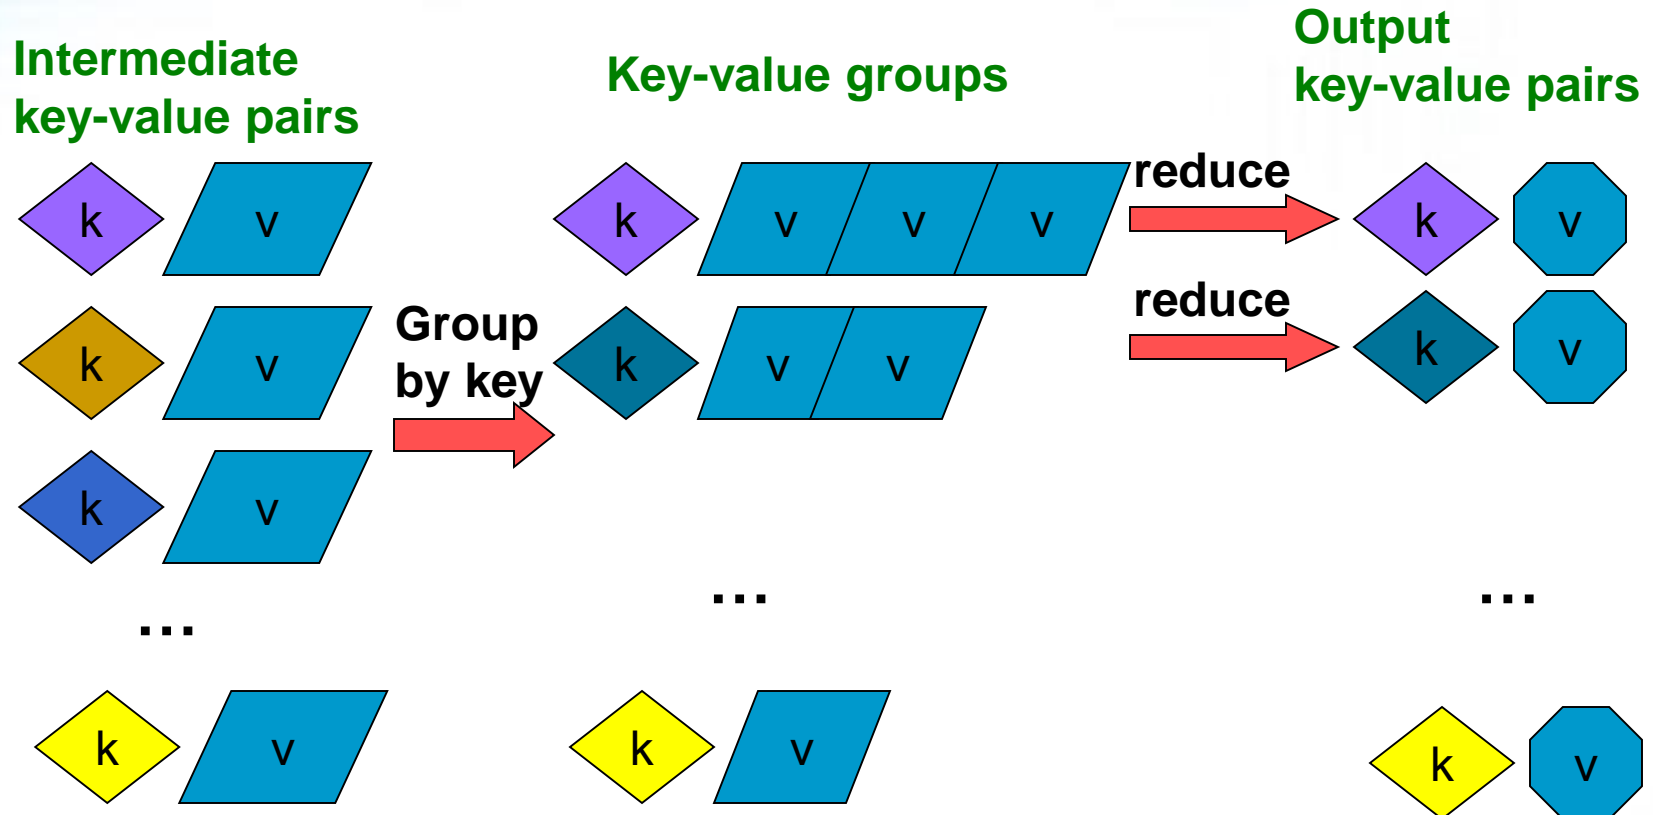
- Sequentially read a lot of data
- Map:
  - Extract something you care about
- Group by key: Sort and Shuffle
- Reduce:
  - Aggregate, summarize, filter or transform
- Write the result

Outline stays the same, **Map** and **Reduce** change to fit the problem

# MapReduce: The Map Step



# MapReduce: The Reduce Step





# More Specifically

- **Input:** a set of key–value pairs
- Programmer specifies two methods:
  - **Map( $k, v$ )**  $\rightarrow \langle k', v' \rangle^*$ 
    - Takes a key–value pair and outputs a set of key–value pairs
      - E.g., key is the filename, value is a single line in the file
    - There is one Map call for every  $(k, v)$  pair
  - **Reduce( $k', \langle v' \rangle^*$ )**  $\rightarrow \langle k', v'' \rangle^*$ 
    - All values  $v'$  with same key  $k'$  are reduced together and processed in  $v'$  order
    - There is one Reduce function call per unique key  $k'$

# MapReduce: Word Counting

Provided by the programmer

## MAP:

Read input and produces a set of key-value pairs

## Group by key:

Collect all pairs with same key

Provided by the programmer

## Reduce:

Collect all values belonging to the key and output

The crew of the space shuttle Endeavor recently returned to Earth as ambassadors, harbingers of a new era of space exploration. Scientists at NASA are saying that the recent assembly of the Dextre bot is the first step in a long term space-based man/machine partnership. "The work we're doing now -- the robotics we're doing -- is what we're going to need .....

Big document

(The, 1)  
(crew, 1)  
(of, 1)  
(the, 1)  
(space, 1)  
(shuttle, 1)  
(Endeavor, 1)  
(recently, 1)  
....

(key, value)



(crew, 1)  
(crew, 1)  
(space, 1)  
(the, 1)  
(the, 1)  
(the, 1)  
(shuttle, 1)  
(recently, 1)  
...

(key, value)

(crew, 2)  
(space, 1)  
(the, 3)  
(shuttle, 1)  
(recently, 1)  
...

(key, value)

Only sequential reads

# Word Count Using MapReduce

---

```
map(key, value):
```

```
// key: document name; value: text of the document
```

```
  for each word w in value:
```

```
    emit(w, 1)
```

```
reduce(key, values):
```

```
// key: a word; value: an iterator over counts
```

```
  result = 0
```

```
  for each count v in values:
```

```
    result += v
```

```
  emit(key, result)
```

# MapReduce: Environment

---

MapReduce environment takes care of:

- Partitioning the input data
- Scheduling the program's execution across a set of machines
- Performing the group by key step
- Handling machine failures
- Managing required inter-machine communication

# MapReduce: A diagram

## MAP:

Read input and produces a set of key-value pairs

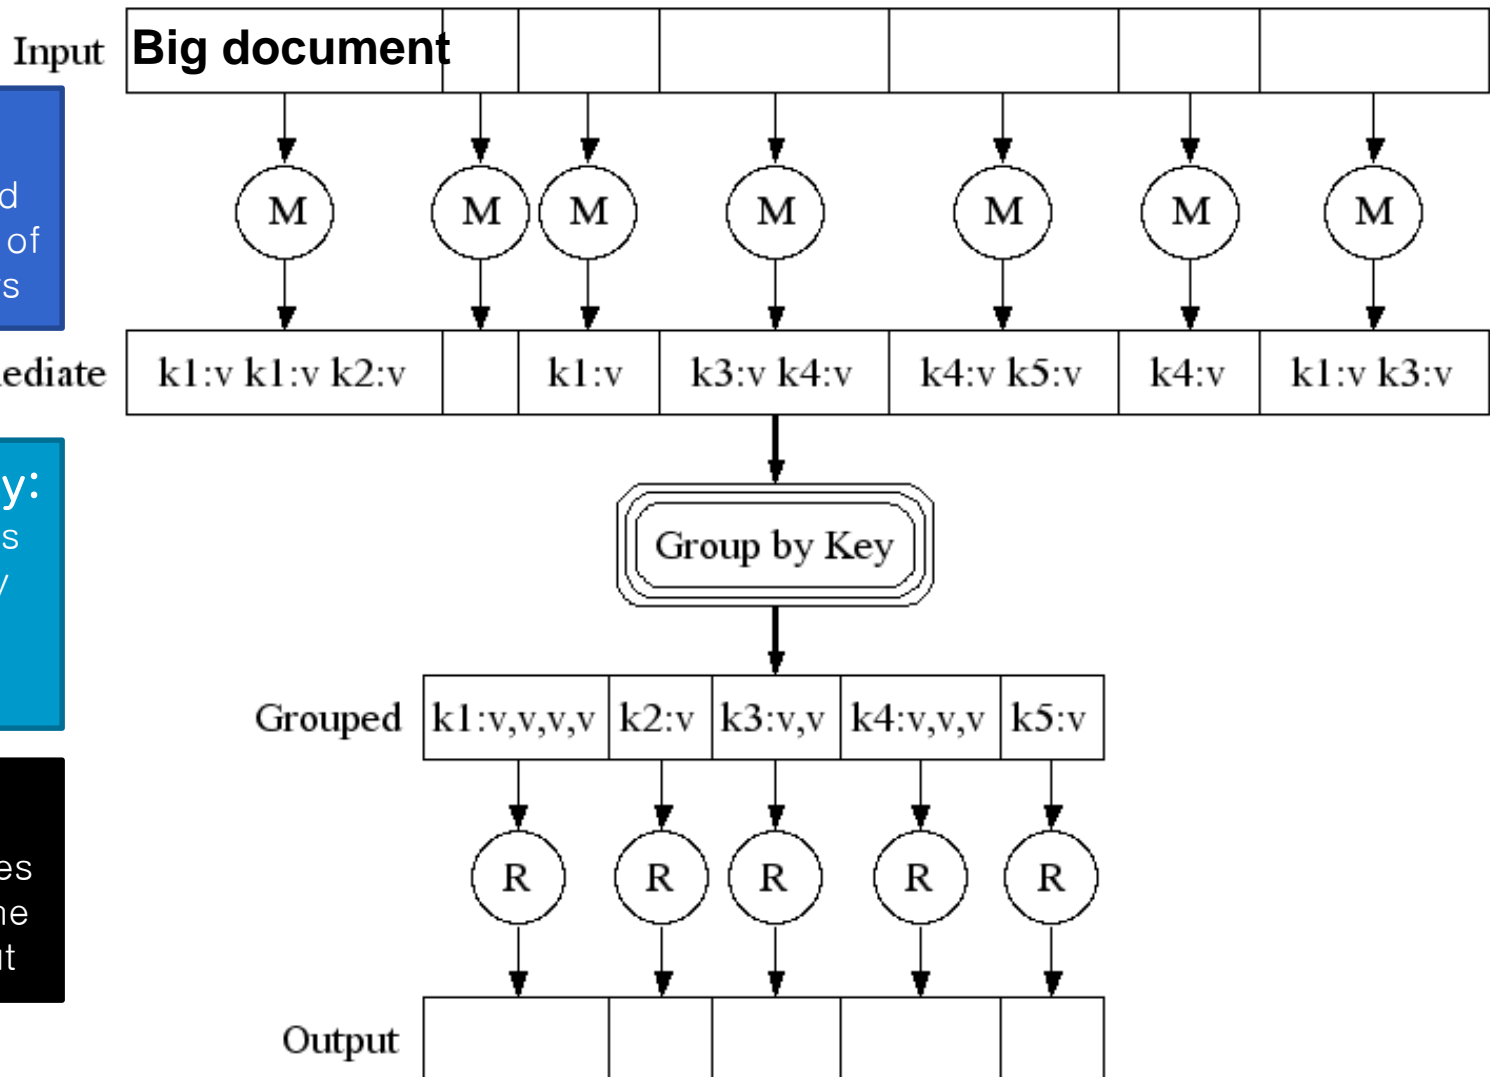
Intermediate

## Group by key:

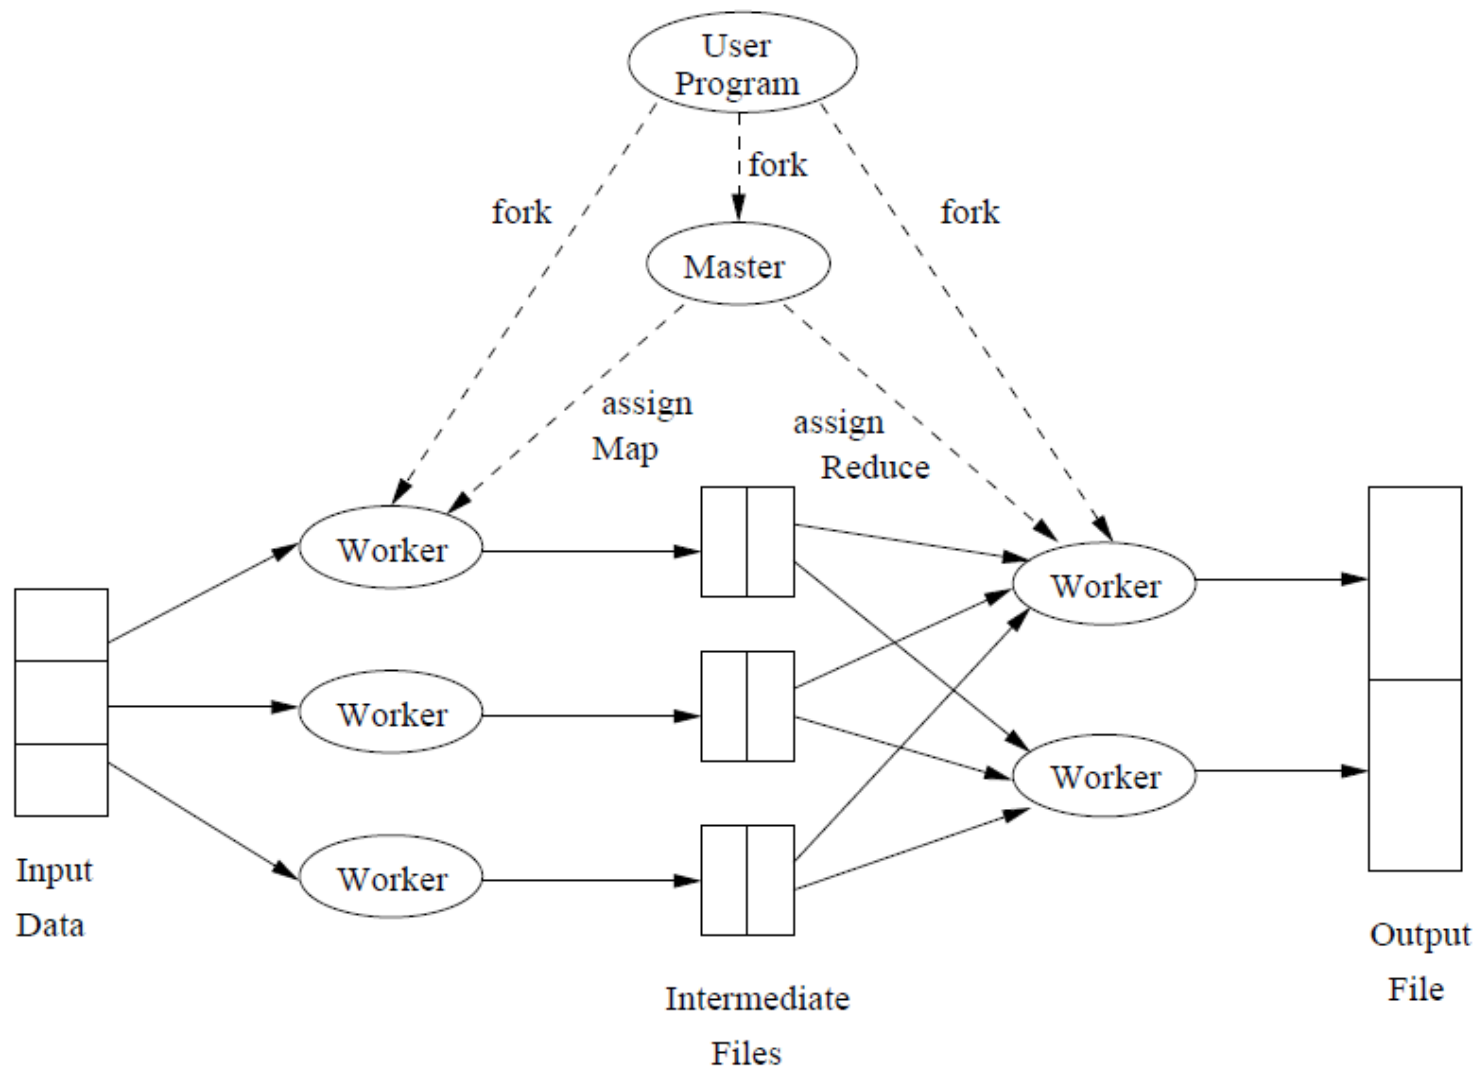
Collect all pairs with same key (Hash merge, Shuffle, Sort, Partition)

## Reduce:

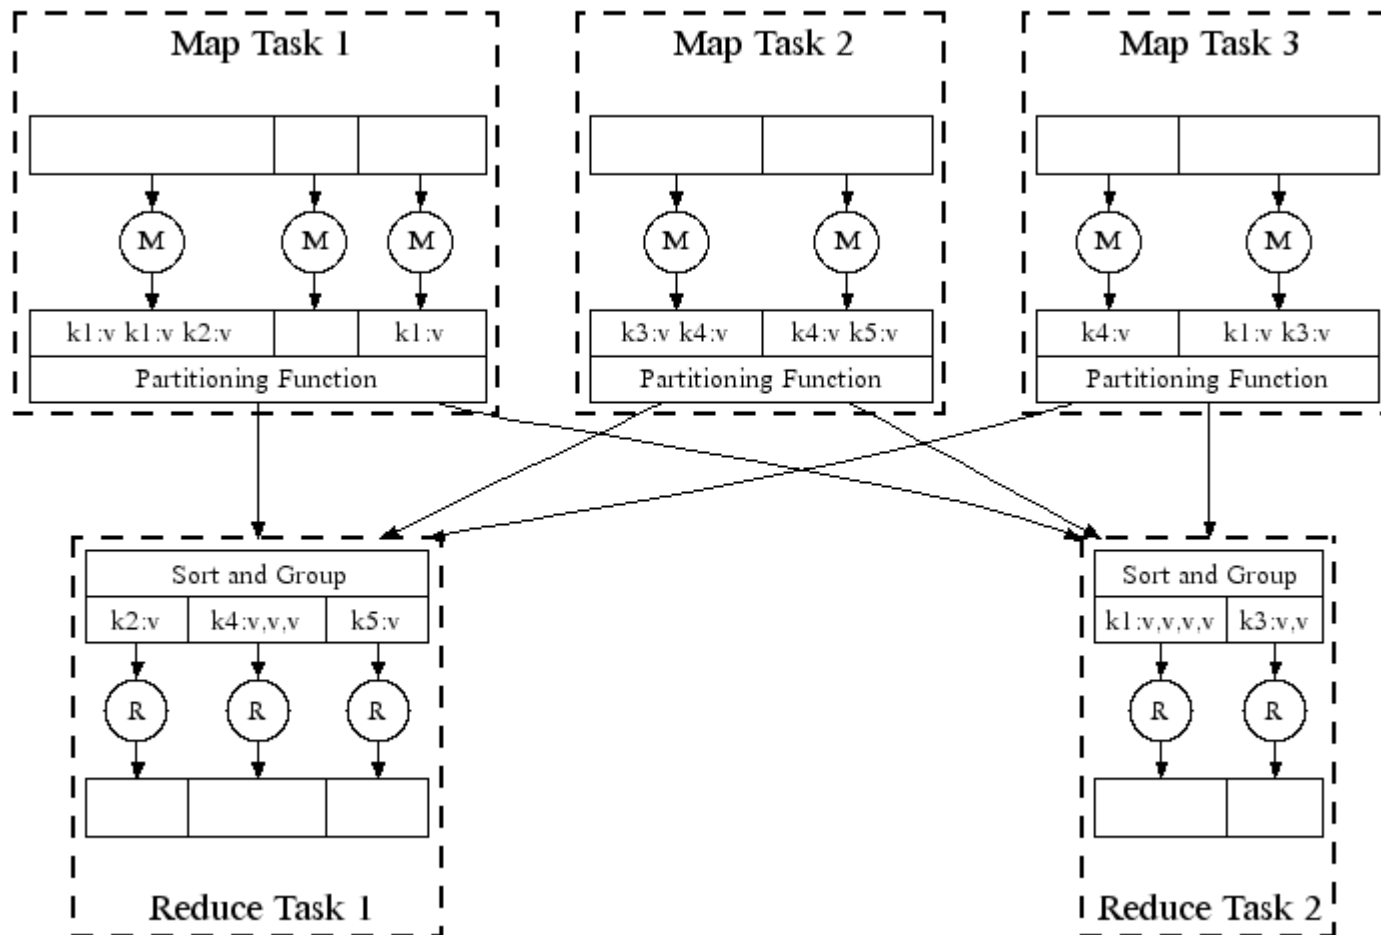
Collect all values belonging to the key and output



# Overview of the execution of a MapReduce program



# MapReduce: In Parallel





# MapReduce

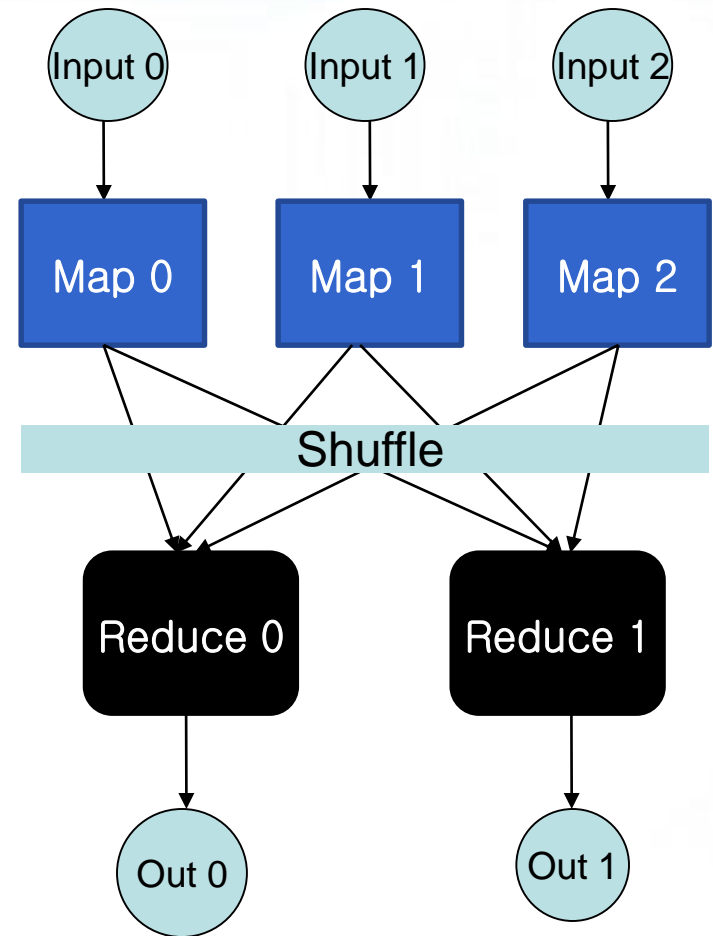
## ■ Programmer specifies:

- Map and Reduce and input files

## ■ Workflow:

- Read inputs as a set of key-value-pairs
- **Map** transforms input kv-pairs into a new set of  $k'v'$ -pairs
- Sorts & Shuffles the  $k'v'$ -pairs to output nodes
- All  $k'v'$ -pairs with a given  $k'$  are sent to the same **reduce**
- **Reduce** processes all  $k'v'$ -pairs grouped by key into new  $k''v''$ -pairs
- Write the resulting pairs to files

## ■ All phases are distributed with many tasks doing the work



# Data Flow

---

- Input and final output are stored on a distributed file system (FS):
  - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- Intermediate results are stored on local FS of Map and Reduce workers
- Output is often input to another MapReduce task

# Coordination: Master

- Master node takes care of coordination:
  - Task status: (idle, in-progress, completed)
  - Idle tasks get scheduled as workers become available
  - When a map task completes, it sends the master the location and sizes of its  $R$  intermediate files, one for each reducer
  - Master pushes this info to reducers
- Master pings workers periodically to detect failures

# Dealing with Failures

---

## ■ Map worker failure

- Map tasks completed or in-progress at worker are reset to idle
- Reduce workers are notified when task is rescheduled on another worker

## ■ Reduce worker failure

- Only in-progress tasks are reset to idle
- Reduce task is restarted

## ■ Master failure

- MapReduce task is aborted and client is notified

# How many Map and Reduce jobs?

- $M$  map tasks,  $R$  reduce tasks

- Rule of a thumb:

- Make  $M$  much larger than the number of nodes in the cluster
- One DFS chunk per map is common
- Improves dynamic load balancing and speeds up recovery from worker failures

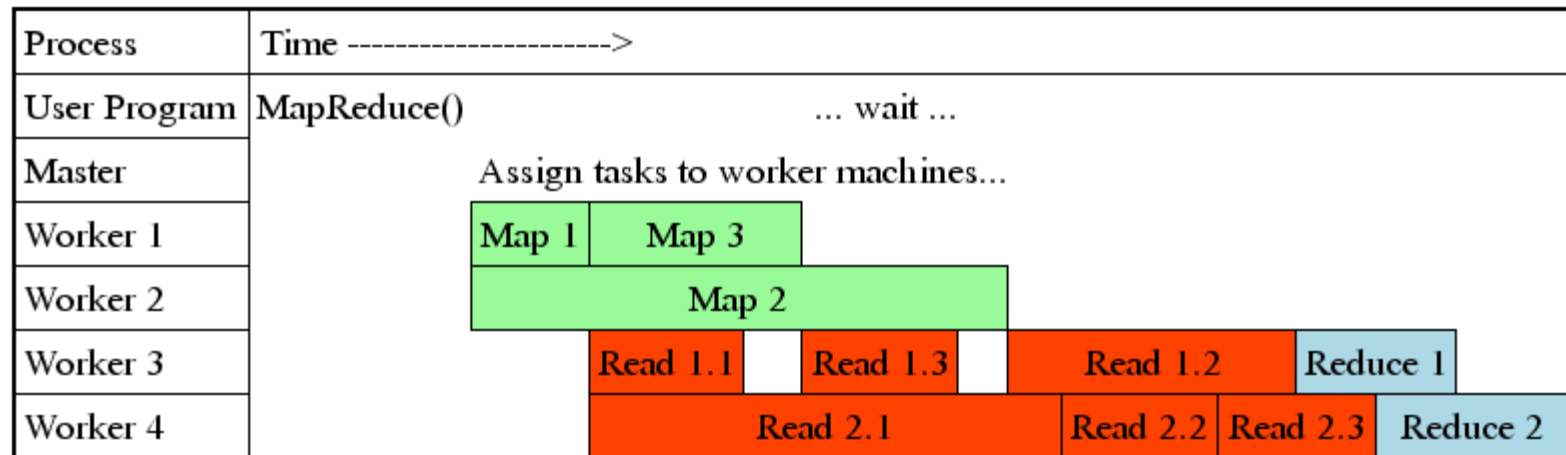
- Usually  $R$  is smaller than  $M$

- Because output is spread across  $R$  files

# Task Granularity & Pipelining

## ■ Fine granularity tasks: map tasks >> machines

- Minimizes time for fault recovery
- Can do pipeline shuffling with map execution
- Better dynamic load balancing



# Refinements: Backup Tasks

## ■ Problem

- Slow workers significantly lengthen the job completion time:
  - Other jobs on the machine
  - Bad disks
  - Weird things

## ■ Solution

- Near end of phase, spawn backup copies of tasks
  - Whichever one finishes first “wins”

## ■ Effect

- Dramatically shortens job completion time

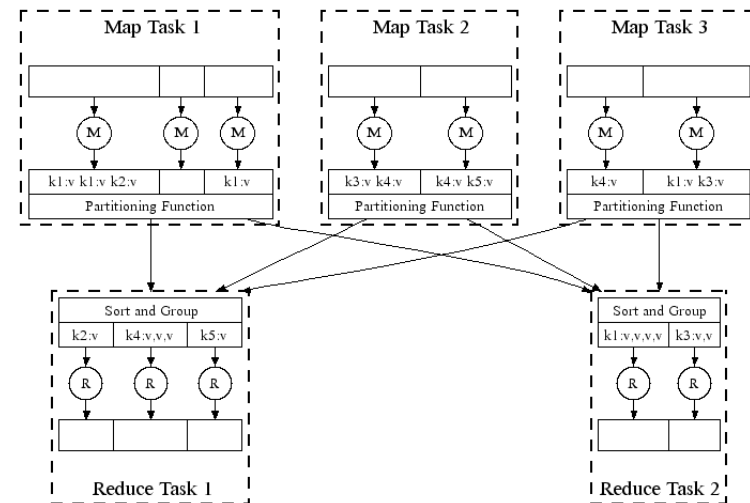


# Refinement: Combiners

- Often a Map task will produce many pairs of the form  $(k, v_1), (k, v_2), \dots$  for the same key  $k$ 
  - E.g., popular words in the word count example
- Can save network time by

- $\text{combine}(k, \text{list}(v_1)) \rightarrow v_2$
- Combiner is usually same as the reduce function

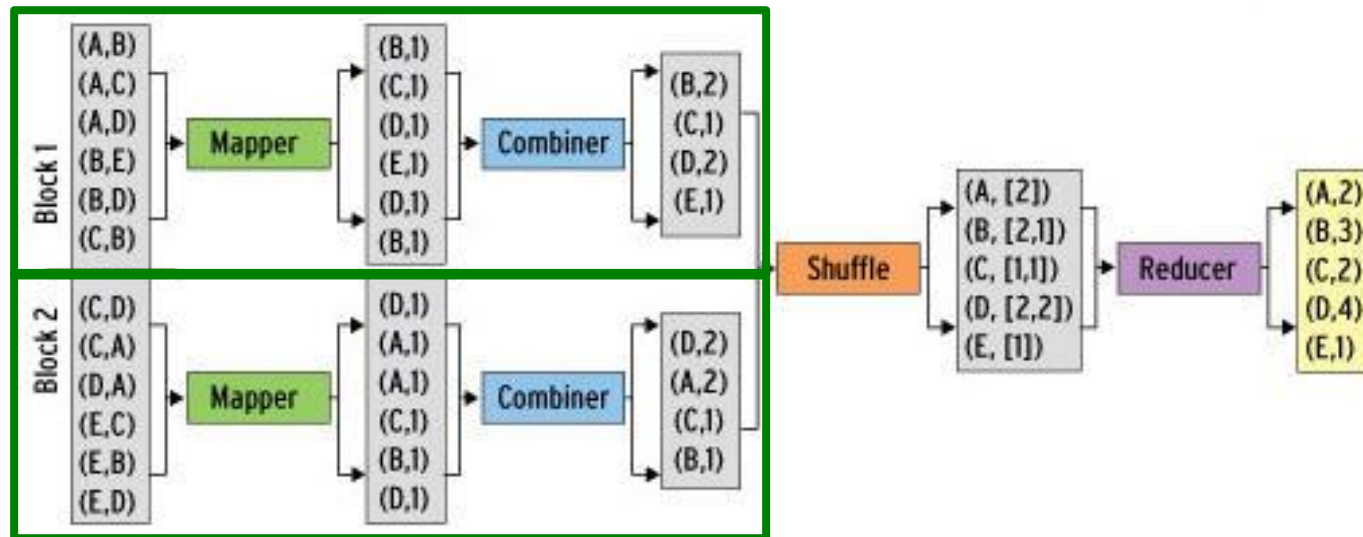
- Works only if reduce function is commutative and associative



# Refinement: Combiners

## ■ Back to our word counting example:

- Combiner combines the values of all keys of a single mapper (single machine):



- Much less data needs to be copied and shuffled!

# Refinement: Partition Function

- Want to control how keys get partitioned
  - Inputs to map tasks are created by contiguous splits of input file
  - Reduce needs to ensure that records with the same intermediate key end up at the same worker
- System uses a default partition function:
  - **$\text{hash}(\text{key}) \bmod R$**
- Sometimes useful to override the hash function:
  - E.g.,  **$\text{hash}(\text{hostname}(\text{URL})) \bmod R$**  ensures URLs from a host end up in the same output file



# Introduction to Apache Hadoop

# Apache Hadoop


- Scalable fault-tolerant distributed system for Big Data: Current version 3.2.1
  - Data Storage
  - Data Processing
  - A virtual Big Data machine
  - Borrowed concepts/Ideas from Google; Open source under the Apache license
  
- Core Hadoop has the following main systems:
  - *Hadoop Common* – contains libraries and utilities needed by other Hadoop modules;
  - *Hadoop Distributed File System (HDFS)* – a distributed file-system that stores data on commodity machines, providing very high aggregate bandwidth across the cluster;
  - *Hadoop YARN* – (introduced in 2012) a platform responsible for managing computing resources in clusters and using them for scheduling users' applications; [\[10\]](#)[\[11\]](#)
  - *Hadoop MapReduce* – an implementation of the MapReduce programming model for large-scale data processing.

# Hadoop 1<sup>st</sup> vs 2<sup>nd</sup>/3<sup>rd</sup>

## Single Use System

Batch Apps

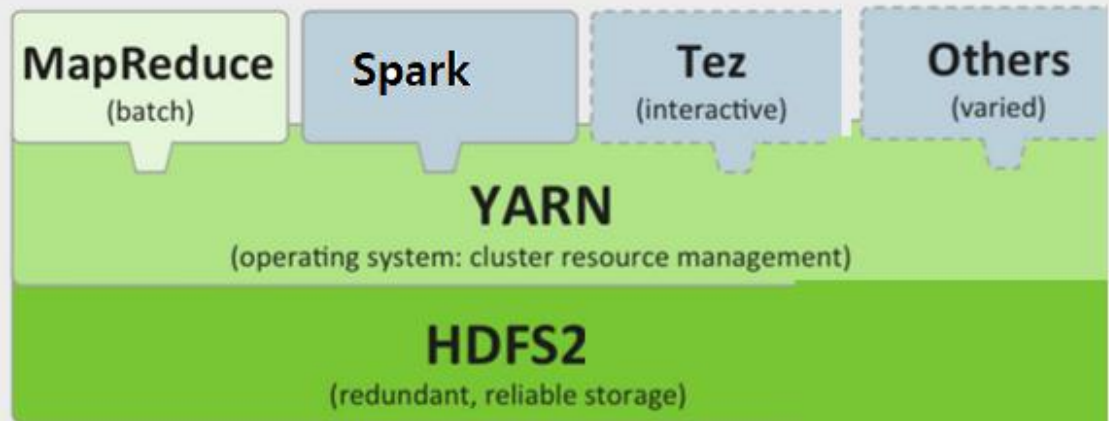
### HADOOP 1.0



## Multi Use Data Platform

Batch, Interactive, Online, Streaming, ...

### HADOOP 2.0/3.0



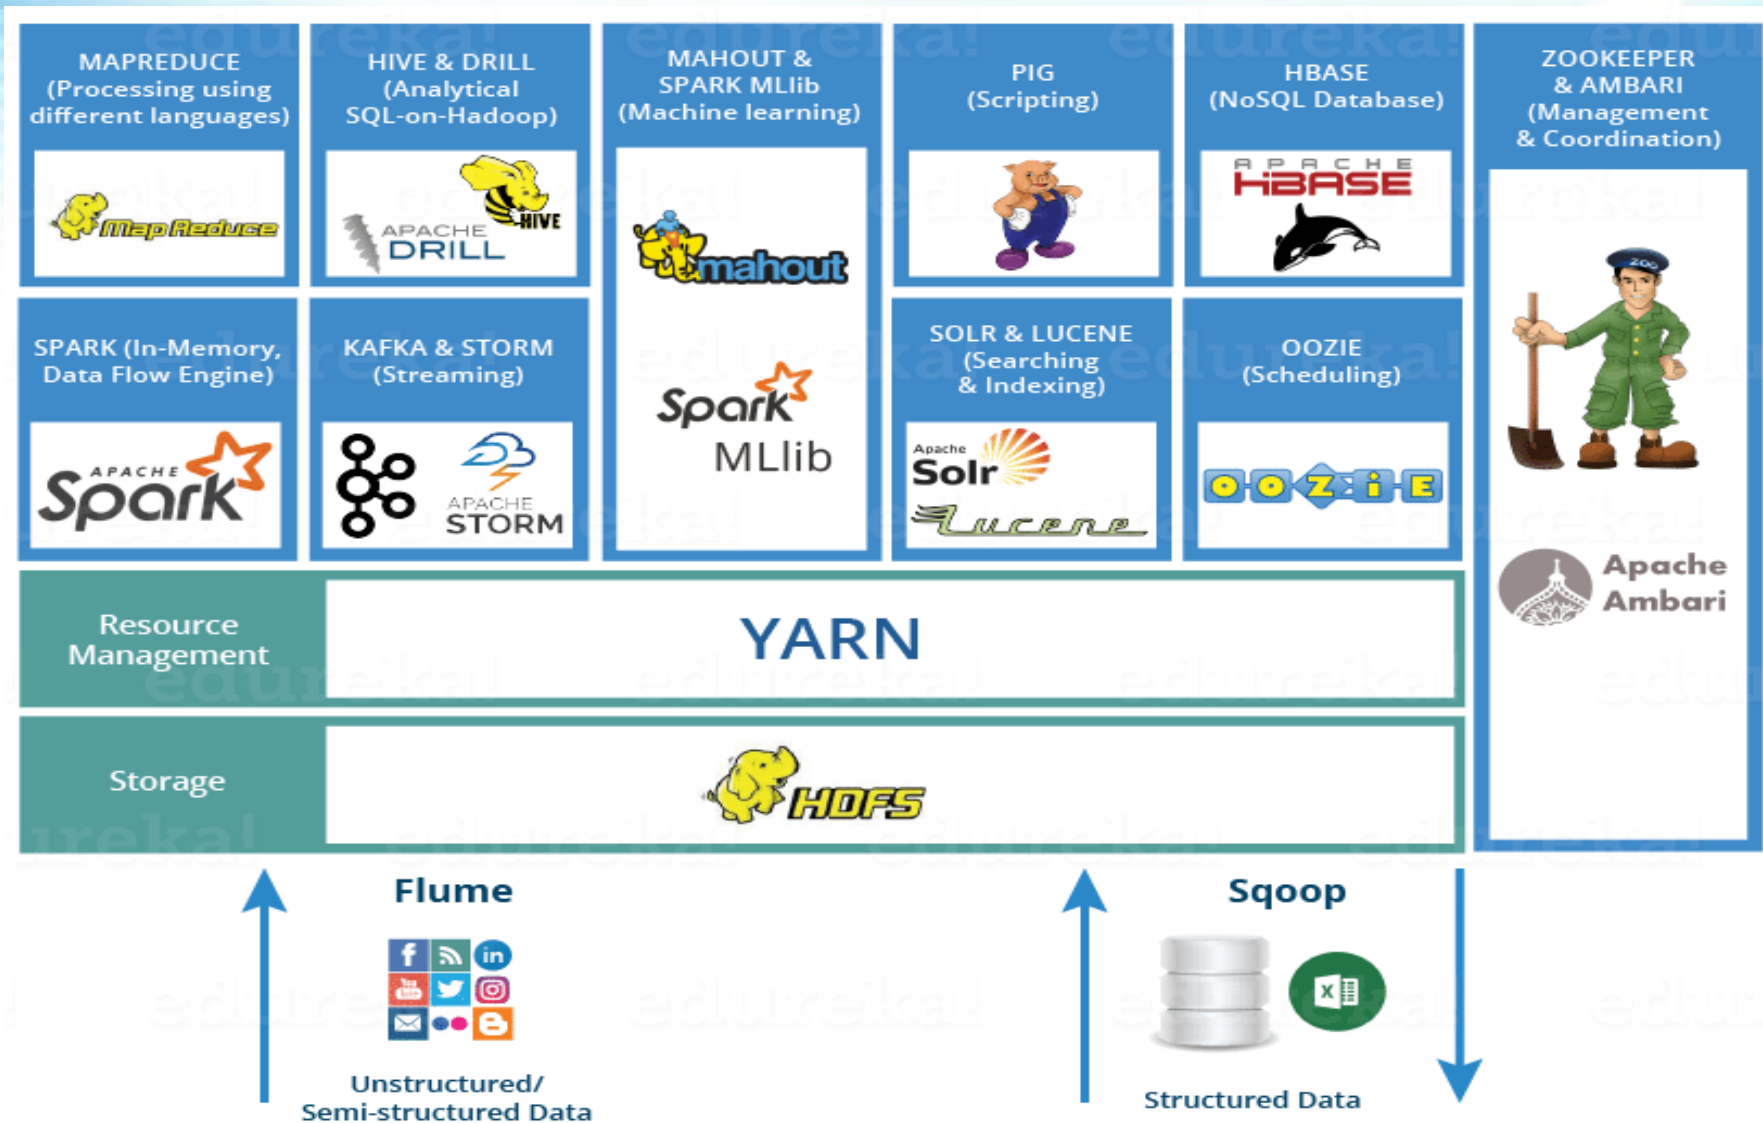


# Hadoop 2.x vs 3.x

Features	Hadoop 2.x	Hadoop 3.x
Min Java Version Required	Java 7	Java 8
Fault Tolerance	Via replication	Via erasure coding
Storage Scheme	3x replication factor for data reliability, 200% overhead	Erasure coding for data reliability, 50% overhead
Yarn Timeline Service	Scalability issues	Highly scalable and reliable
Standby NN	Supports only 1 SBNN	Supports only 2 or more SBNN
Heap Management	We need to configure HADOOP_HEAPSIZE	Provides auto-tuning of heap



# Hadoop Ecosystem (Hadoop 2&3)



# Hadoop: MapReduce

Chapter 2. MapReduce,  
“Hadoop, The Definitive Guide,” 4<sup>th</sup> Ed.,  
Tom White, O’Reilly

# MapReduce Data Flow

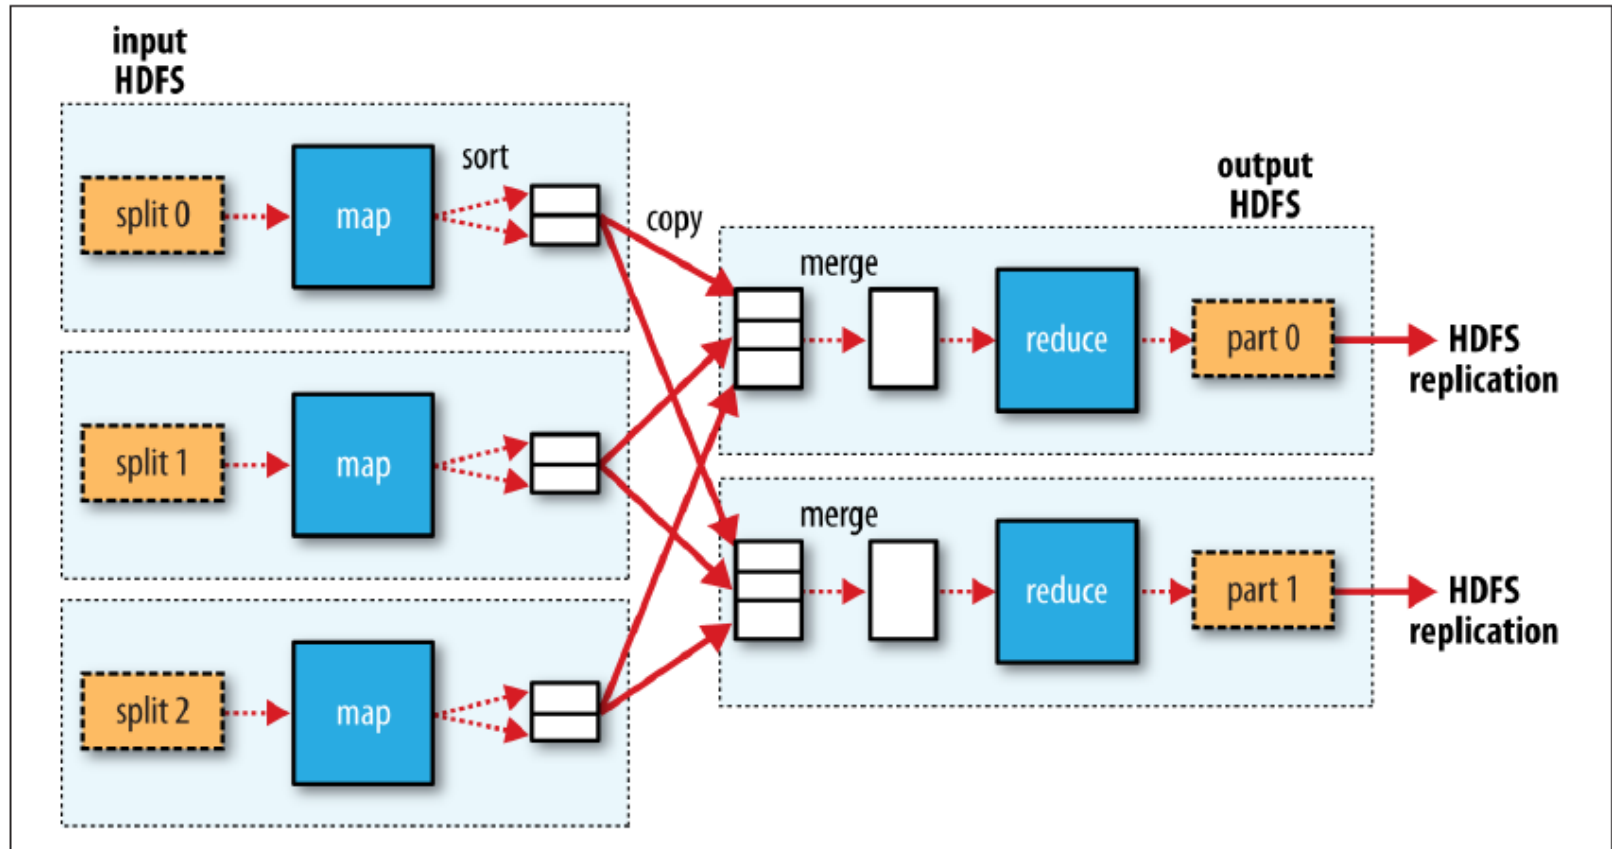


Figure 2-4. MapReduce data flow with multiple reduce tasks

Figure is from Hadoop, The Definitive Guide, 4<sup>th</sup> Edition, Tom White, O'Reilly

# Word Count with Combiner

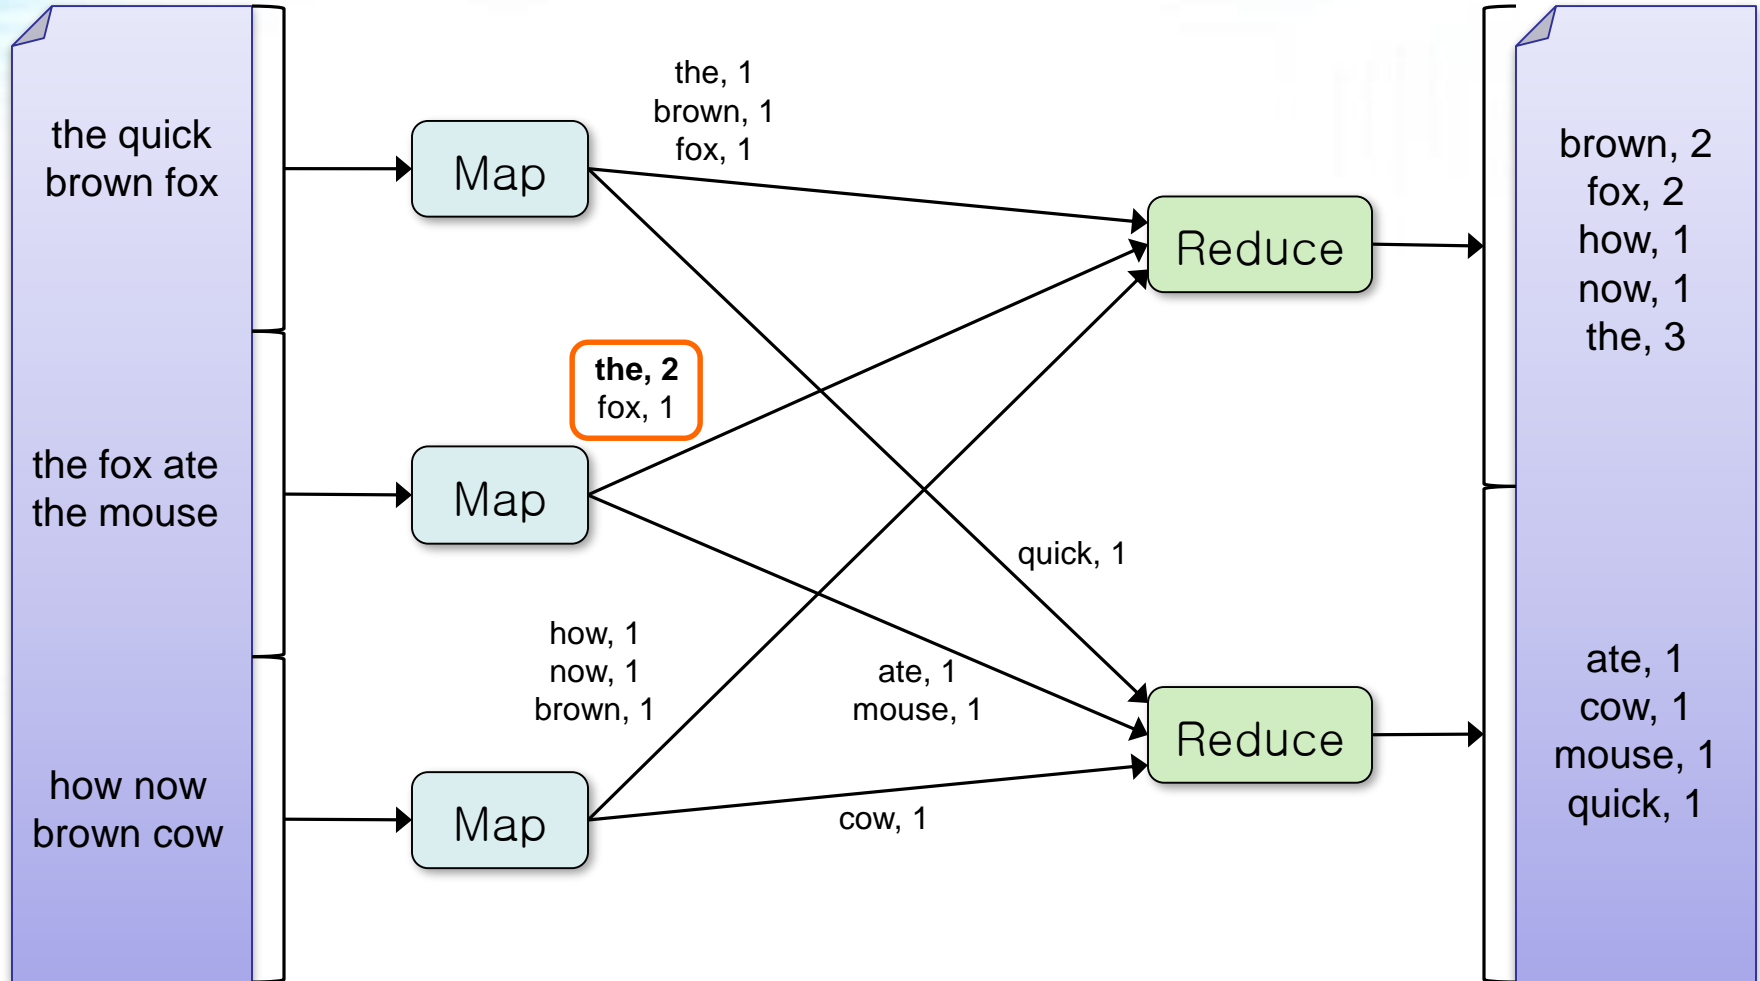
Input

Map &amp; Combine

Shuffle &amp; Sort

Reduce

Output





# HDFS: Hadoop File System

Chapter 3. The Hadoop Distributed Filesystem,  
“Hadoop, The Definitive Guide,” 4<sup>th</sup> Ed., Tom  
White, O’Reilly

# From GFS to HDFS

---

## ■ Terminology differences:

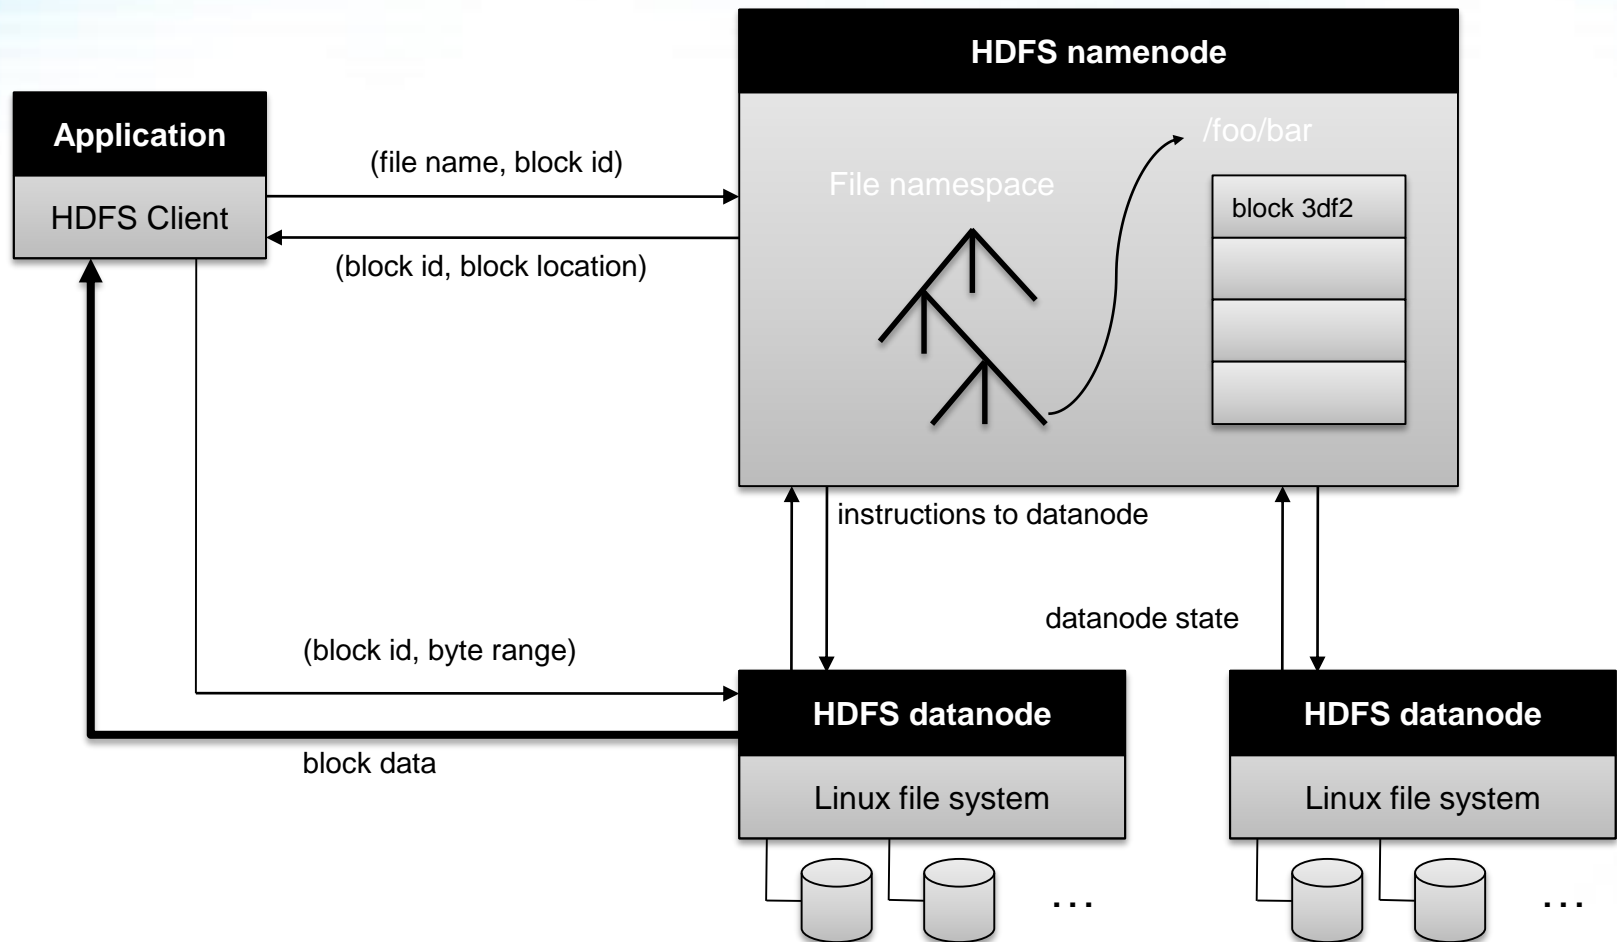
- GFS master = Hadoop namenode
- GFS chunkservers = Hadoop datanodes

## ■ Functional differences:

- HDFS performance is (likely) slower

For the most part, we'll use the Hadoop terminology...

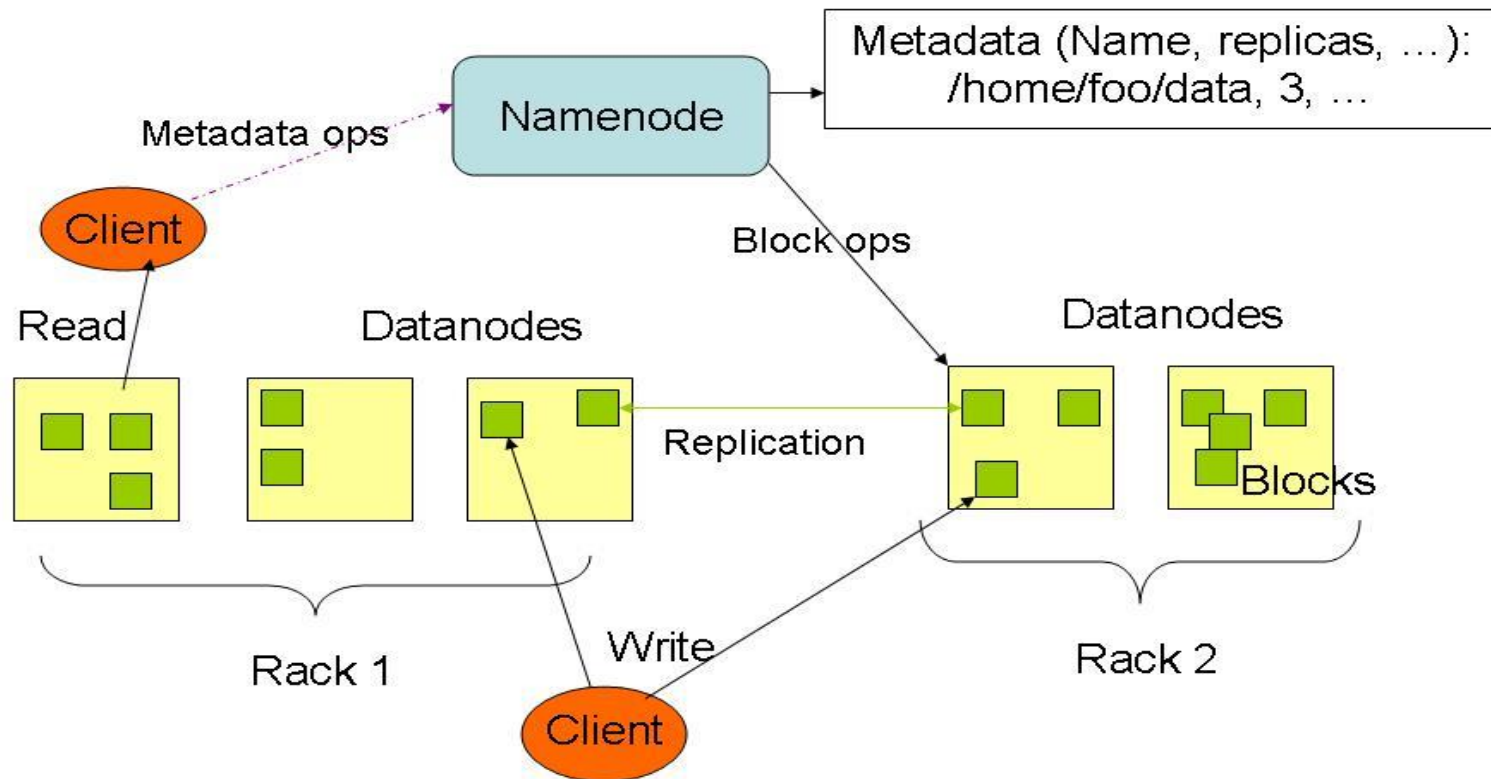
# HDFS Working Flow





# HDFS Architecture

## HDFS Architecture



# Distributed File System

---

- Single Namespace for entire cluster
- Data Coherency
  - Write–once–read–many access model
  - Client can only append to existing files
- Files are broken up into blocks
  - Typically 128 MB block size
  - Each block replicated on multiple DataNodes
- Intelligent Client
  - Client can find location of blocks
  - Client accesses data directly from DataNode

# NameNode Metadata

---

- Meta-data in Memory
  - The entire metadata is in main memory
  - No demand paging of meta-data
- Types of Metadata
  - List of files
  - List of Blocks for each file
  - List of DataNodes for each block
  - File attributes, e.g creation time, replication factor
- A Transaction Log
  - Records file creations, file deletions. etc

# Namenode Responsibilities

---

- **Managing the file system namespace:**
  - Holds file/directory structure, metadata, file-to-block mapping, access permissions, etc.
  
- **Coordinating file operations:**
  - Directs clients to datanodes for reads and writes
  - No data is moved through the namenode
  
- **Maintaining overall health:**
  - Periodic communication with the datanodes
  - Block re-replication and rebalancing
  - Garbage collection

# DataNode

---

- A Block Server
  - Stores data in the local file system (e.g. ext3)
  - Stores meta-data of a block (e.g. CRC)
  - Serves data and meta-data to Clients
  
- Block Report
  - Periodically sends a report of all existing blocks to the NameNode
  
- Facilitates Pipelining of Data
  - Forwards data to other specified DataNodes

# Block Placement

---

- Current Strategy
  - One replica on local node
  - Second replica on a remote rack
  - Third replica on same remote rack
  - Additional replicas are randomly placed
- Clients read from nearest replica
- Would like to make this policy pluggable

# Data Correctness

---

- Use Checksums to validate data
  - Use CRC32
- File Creation
  - Client computes checksum per 512 byte
  - DataNode stores the checksum
- File access
  - Client retrieves the data and checksum from DataNode
  - If Validation fails, Client tries other replicas



# NameNode Failure

---

- A single point of failure
- Transaction Log stored in multiple directories
  - A directory on the local file system
  - A directory on a remote file system (NFS/CIFS)
- Need to develop a real HA(High Availability) solution

# Anatomy of a File Read in HDFS

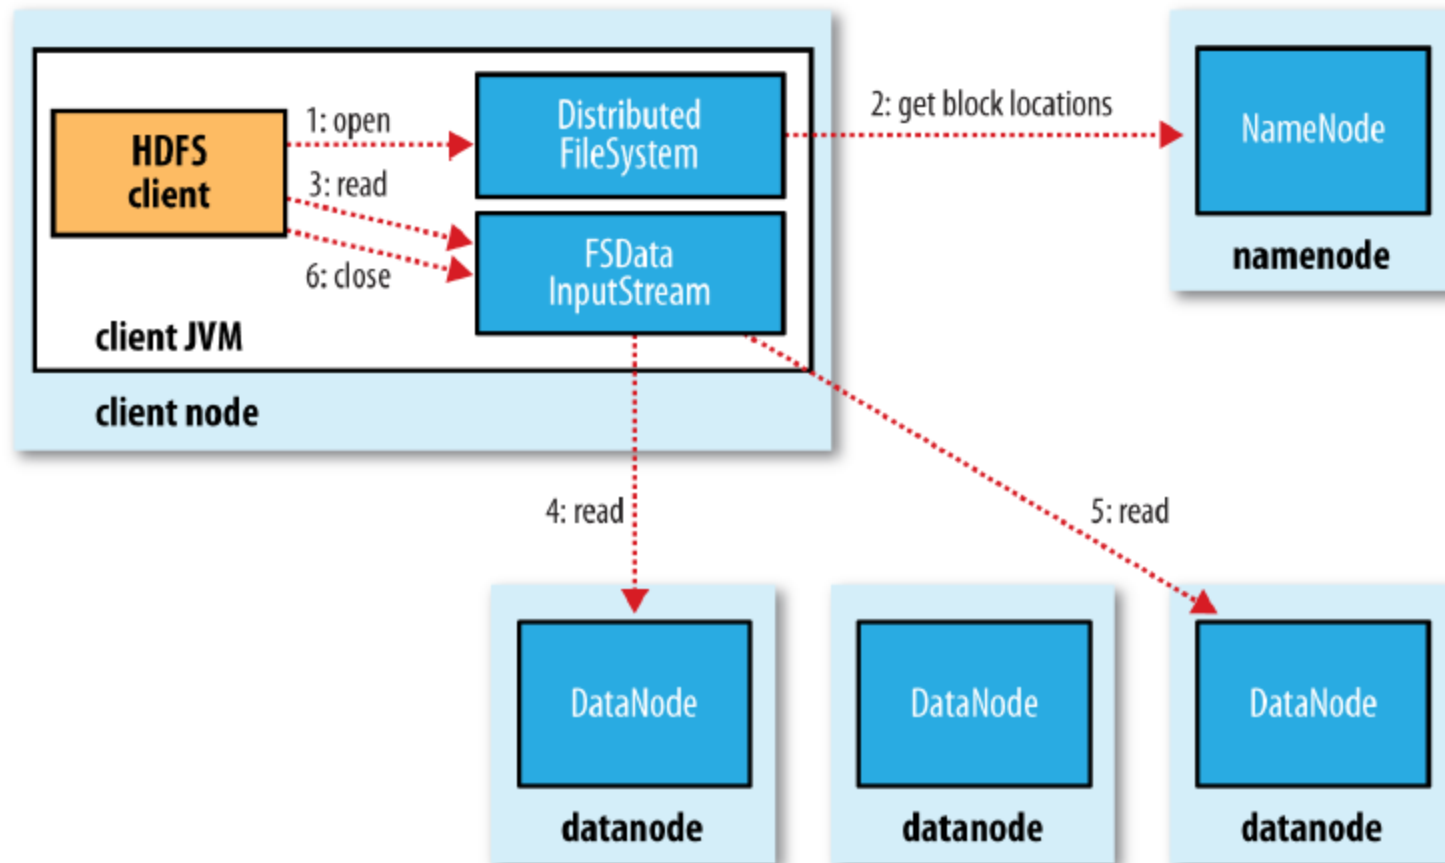


Figure 3-2. A client reading data from HDFS

Figure is from Hadoop, The Definitive Guide, 4<sup>th</sup> Edition, Tom White, O'Reilly

# Network Distance in HDFS

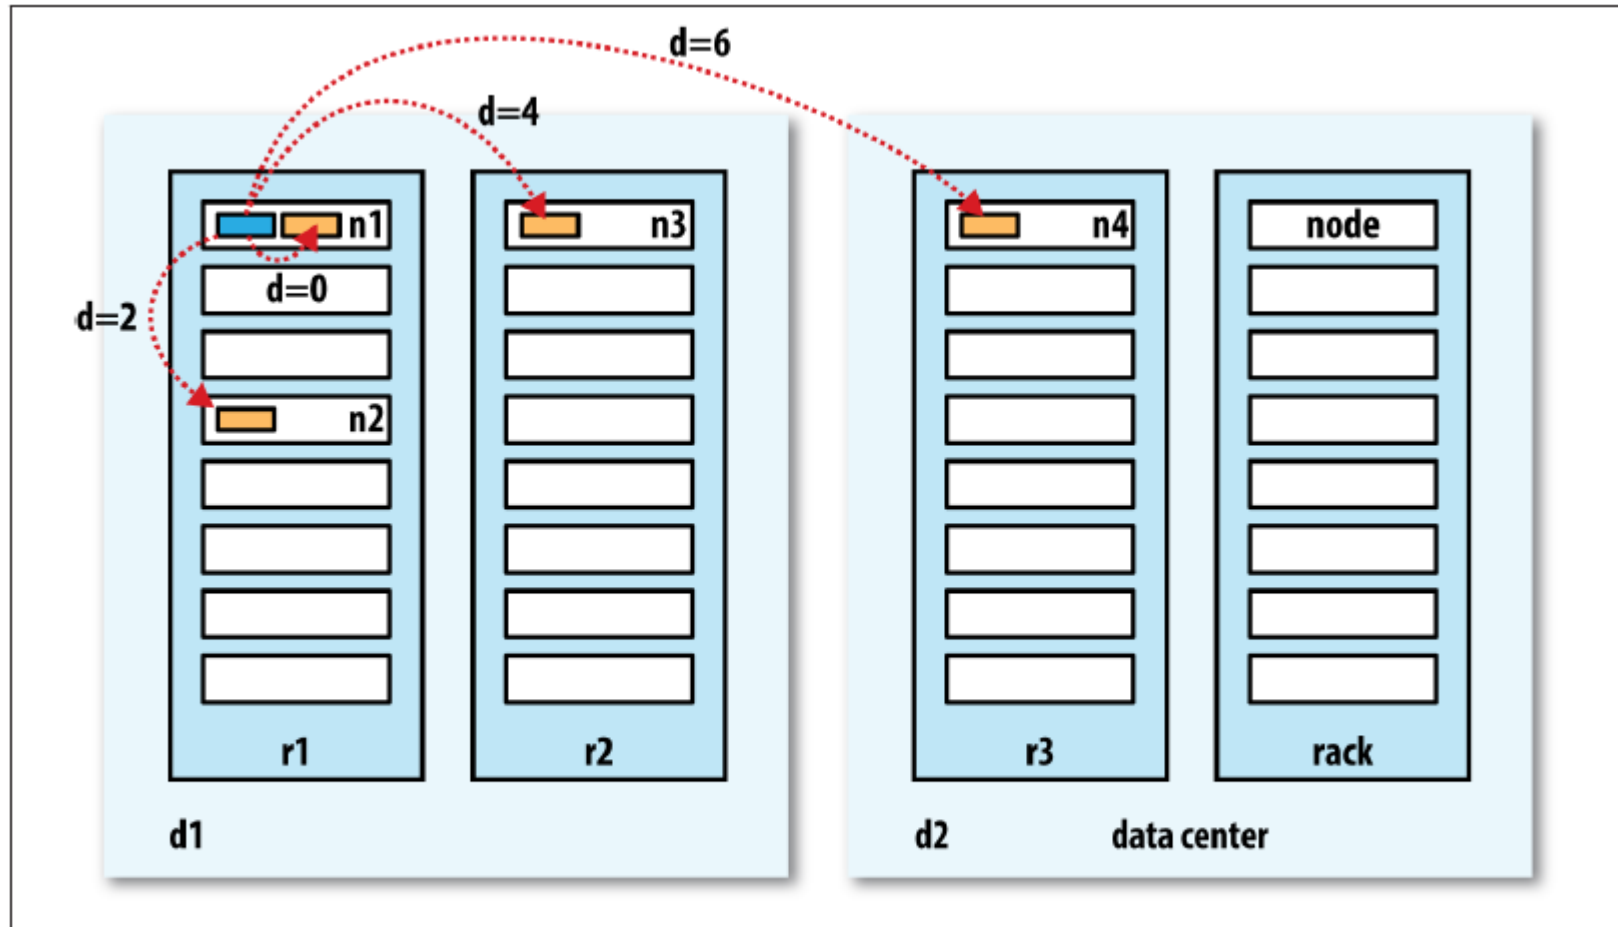


Figure 3-3. Network distance in Hadoop

Figure is from Hadoop, The Definitive Guide, 4<sup>th</sup> Edition, Tom White, O'Reilly

# Anatomy of a File Write in HDFS

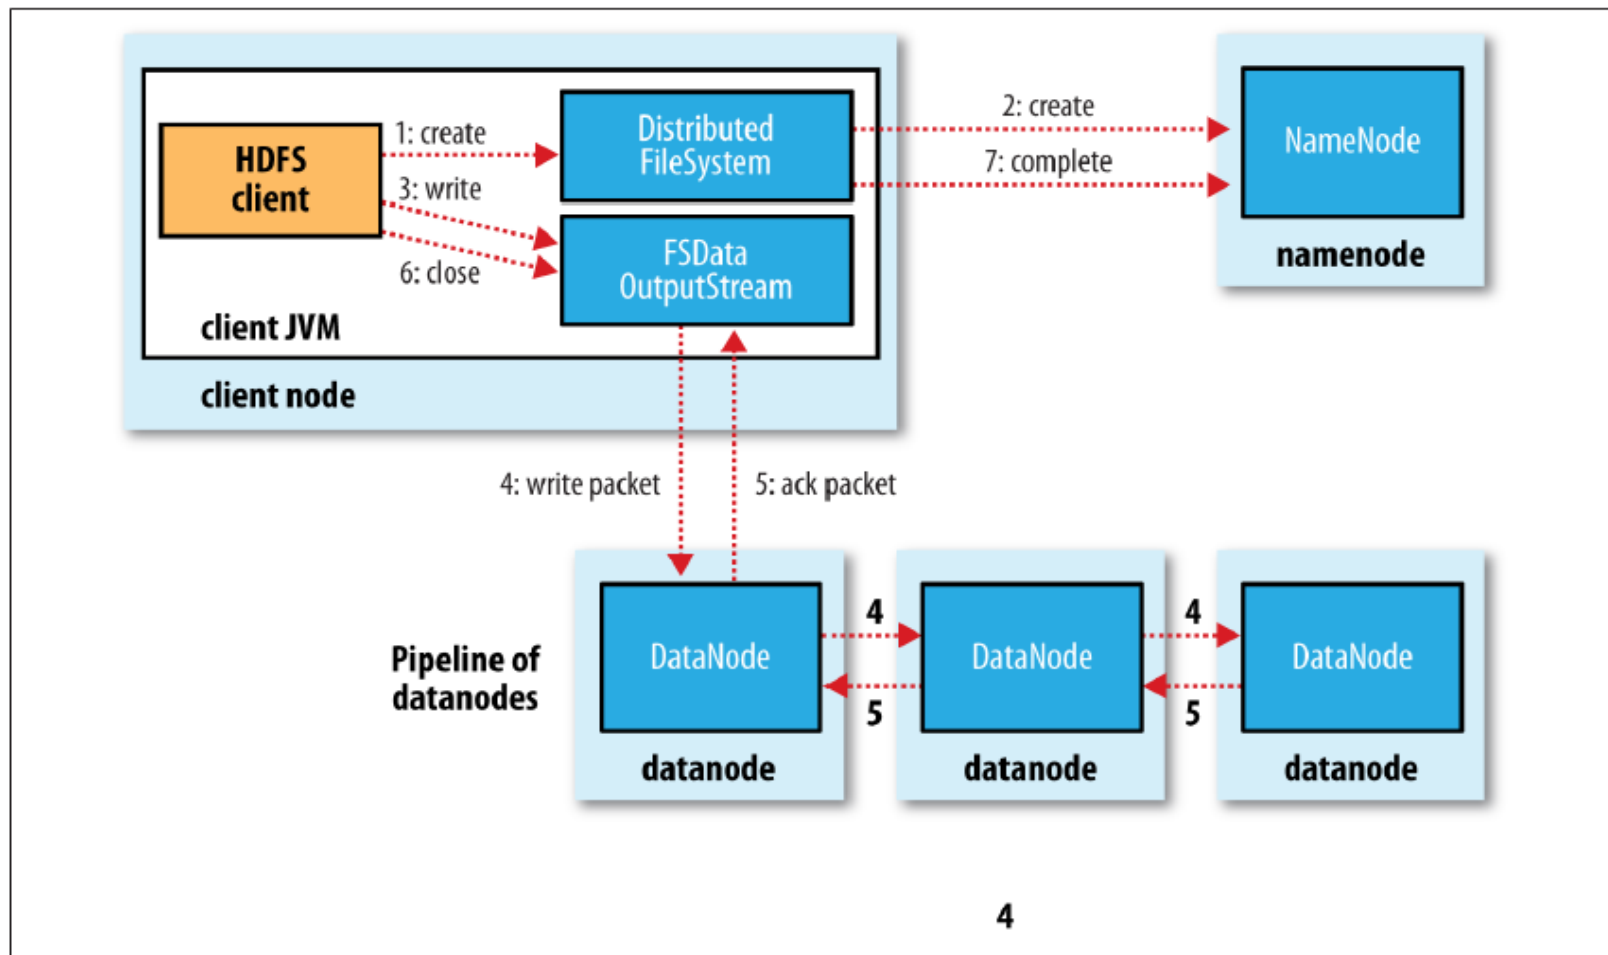


Figure 3-4. A client writing data to HDFS

Figure is from Hadoop, The Definitive Guide, 4<sup>th</sup> Edition, Tom White, O'Reilly

# Replica Placement Strategies in Hadoop (Replicas = 3)

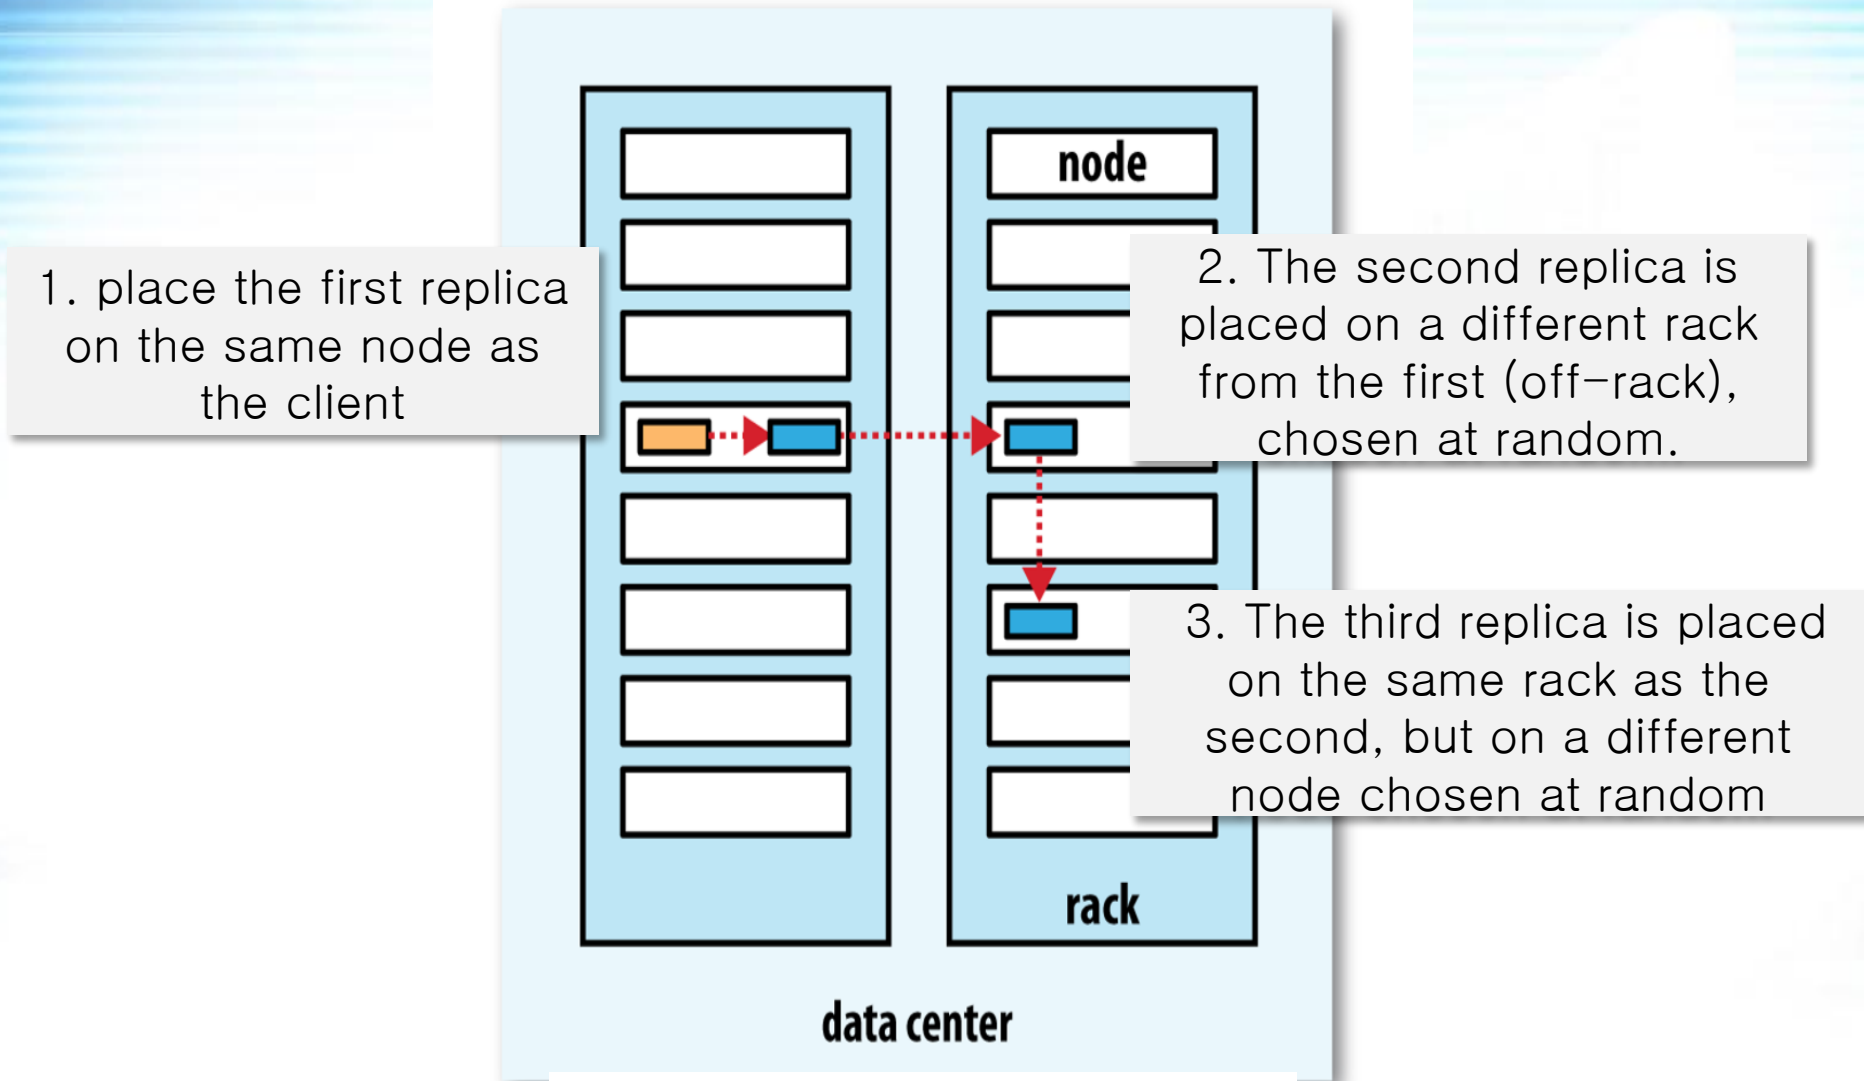


Figure 3-5. A typical replica pipeline



# Apache Hadoop YARN: Yet Another Resource Negotiator

Chapter 4. YARN,  
“Hadoop, The Definitive  
Guide,” 4<sup>th</sup> Ed., Tom White,  
O’Reilly

# YARN

---

- Scalability
- Multi-tenancy
- Serviceability
- Locality Awareness
- High Cluster Utilization
- Reliability/Availability
- Secure and auditable operation
- Support for Programming Model Diversity

## ■ Flexible Resource Model

- Hadoop: # of Map/reduce slots are fixed.
- Easy, but lower utilization



# YARN

---

- Scalability
- Multi-tenancy
- Serviceability
- Locality Awareness
- High Cluster Utilization
- Reliability/Availability
- Secure and auditable operation
- Support for Programming Model Diversity
- Flexible Resource Model

## ■ Backward Compatibility

- The system behaves similar to the old Hadoop

# YARN

---

- Separating resource management functions from the programming model
- MapReduce becomes just one of the application
- Dryad, ... Etc
- Binary compatible/Source compatible

# YARN: Architecture

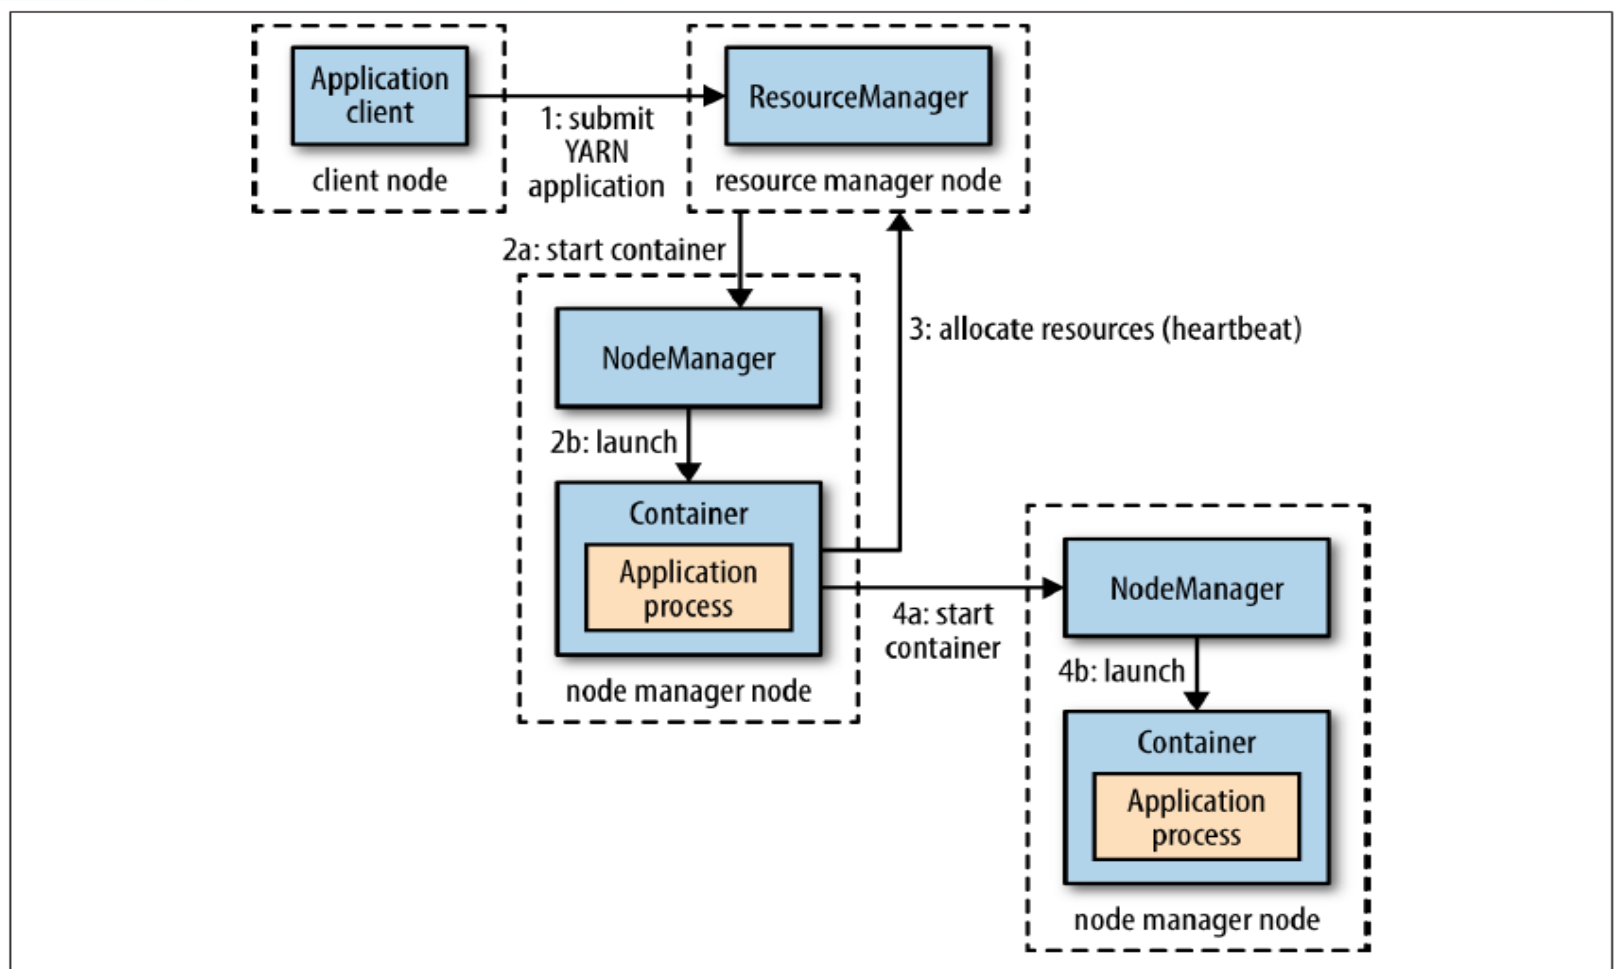


Figure 4-2. How YARN runs an application

Figure is from Hadoop, The Definitive Guide, 4<sup>th</sup> Edition, Tom White, O'Reilly

# Resource Manager

---

- One per cluster

- Central, global view
- Enable global properties
  - Fairness, capacity, locality

- Container

- Logical bundle of resources (CPU/memory)

- Job requests are submitted to RM

- To start a job, RM finds a container to spawn AM

- No static resource partitioning

# Resource Manager (cont' )

---

- only handles an overall resource profile for each application
  - Local optimization/internal flow is up to the application
  
- Preemption
  - Request resources back from an application
  - Checkpoint snapshot instead of explicitly killing jobs / migrate computation to other containers

# Application Master

---

- The head of a job
- Runs as a container
- Request resources from RM
  - # of containers/ resource per container/ locality ...
- Dynamically changing resource consumption
- Can run any user code (Dryad, MapReduce, Tez, REEF...etc)
- Requests are “*late-binding*”

# MapReduce AM

---

- Optimizes for locality among map tasks with identical resource requirements
  - Selecting a task with input data close to the container.
- AM determines the semantics of the success or failure of the container

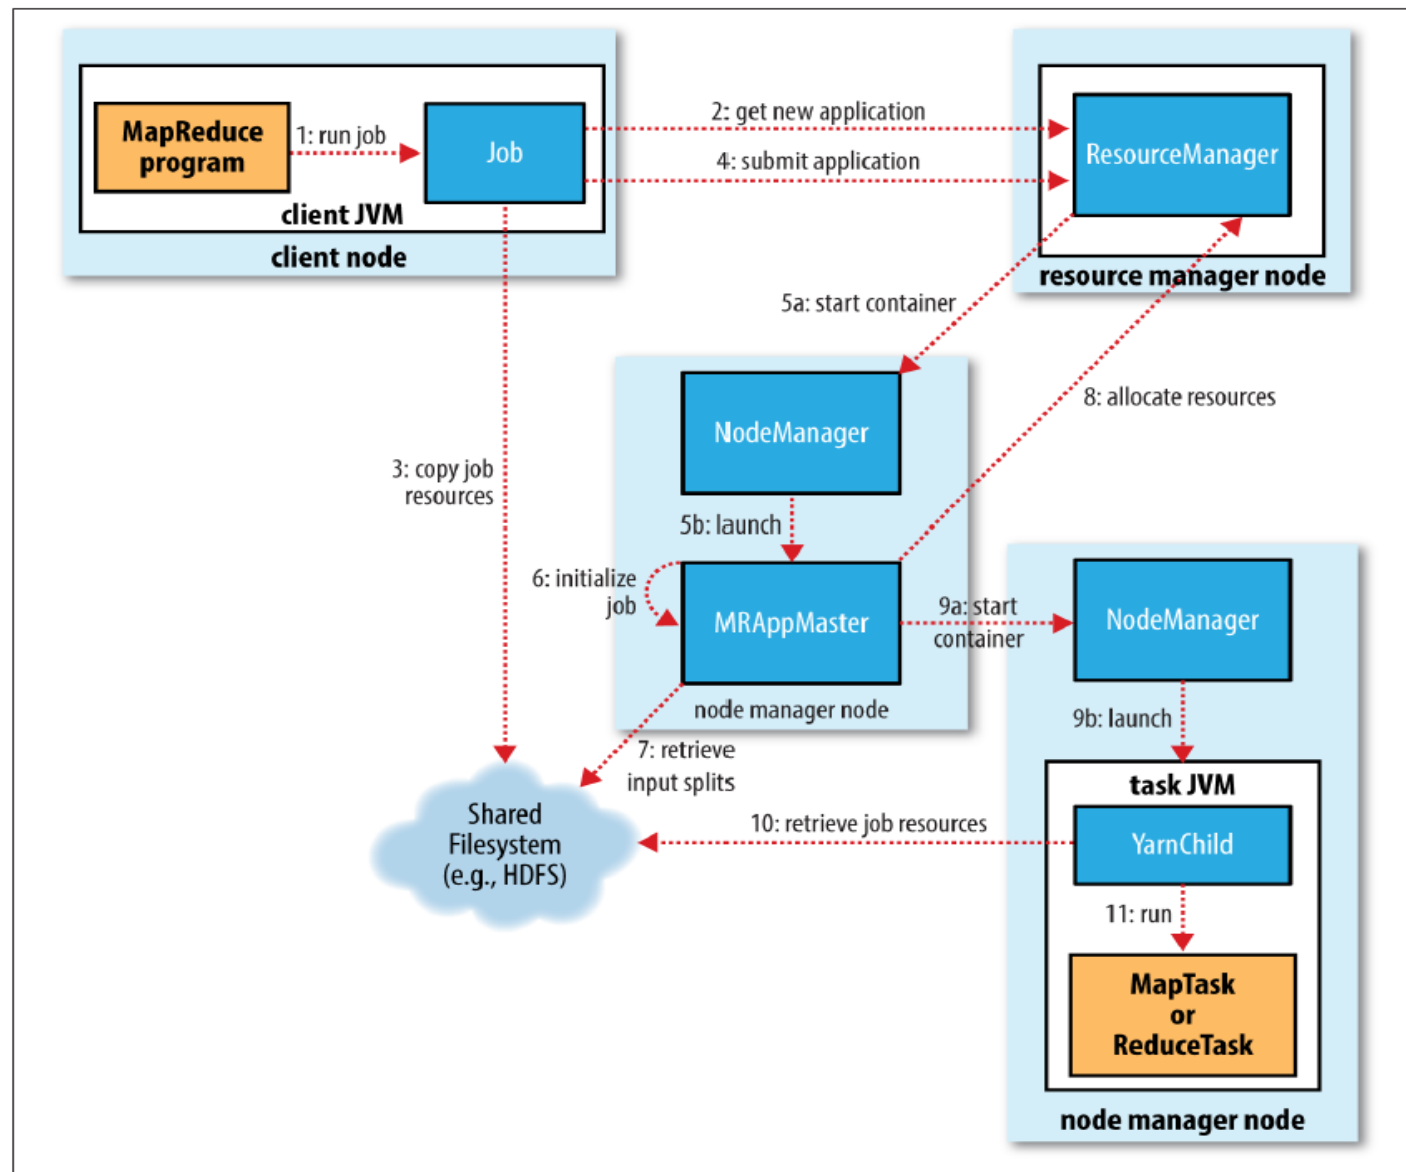


# Node Manager

---

- The “worker” daemon. Registers with RM
- One per node
- Container Launch Context – env var, commands...
- Report resources (memory/CPU/etc...)
- Configure the environment for task execution
- Garbage collection/ Authentication
- Auxiliary services
  - Output intermediate data between map and reduce tasks

# MapReduce Job in Hadoop YARN



# Fault tolerance and availability

---

## ■ RM Failure

- Recover using persistent storage
- Kill all containers, including AMs'
- Relaunch AMs

## ■ NM Failure

- RM detects it, mark the containers as killed, report to Ams

## ■ AM Failure

- RM kills the container and restarts it.

## ■ Container Failure

- The framework is responsible for recovery

# Hadoop Resources



# Reading

---

- Jeffrey Dean and Sanjay Ghemawat:  
MapReduce: Simplified Data Processing on  
Large Clusters

➤ <http://labs.google.com/papers/mapreduce.html>

- Sanjay Ghemawat, Howard Gobioff, and  
Shun-Tak Leung: The Google File System

➤ <http://labs.google.com/papers/gfs.html>



# Resources

---

## ■ Hadoop Wiki

- Introduction
  - <http://wiki.apache.org/lucene-hadoop/>
- Getting Started
  - <http://wiki.apache.org/lucene-hadoop/GettingStartedWithHadoop>
- Map/Reduce Overview
  - <http://wiki.apache.org/lucene-hadoop/HadoopMapReduce>
  - <http://wiki.apache.org/lucene-hadoop/HadoopMapReduceClasses>
- Eclipse Environment
  - <http://wiki.apache.org/lucene-hadoop/EclipseEnvironment>

## ■ Javadoc

- <http://lucene.apache.org/hadoop/docs/api/>

# Resources

---

## ■ Releases from Apache download mirrors

➤ <http://www.apache.org/dyn/closer.cgi/lucene/hadoop/>

## ■ Nightly builds of source

➤ <http://people.apache.org/dist/lucene/hadoop/nightly/>

## ■ Source code from subversion

➤ [http://lucene.apache.org/hadoop/version\\_control.html](http://lucene.apache.org/hadoop/version_control.html)



# Introduction to Apache Spark



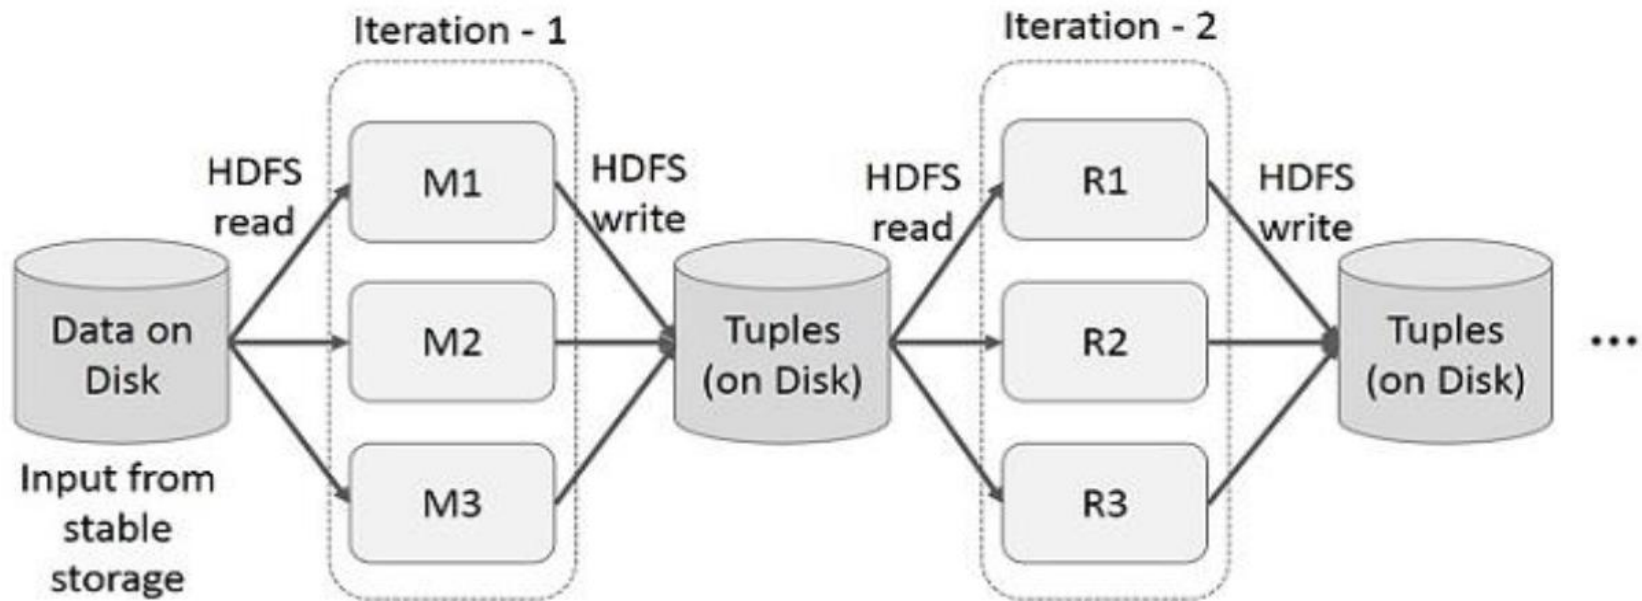
# Shortcoming of MapReduce

---

- Forces your data processing into Map and Reduce
  - Other workflows missing include join, filter, flatMap, groupByKey, union, intersection, ...
- Based on “Acyclic Data Flow” from Disk to Disk (HDFS)
- Read and write to Disk before and after Map and Reduce (stateless machine)
  - Not efficient for iterative tasks, i.e. Machine Learning
- Only for Batch processing
  - Interactivity, streaming data

# Iteration in MapReduce

- Both Iterative and Interactive applications require faster data sharing across parallel jobs.
- Data sharing is slow in MapReduce due to replication, serialization, and disk IO. Regarding storage system, most of the Hadoop applications, they spend more than 90% of the time doing HDFS read-write operations.




**Figure:** Iterative operations on MapReduce



# Interactiveness in MapReduce

- Both Iterative and Interactive applications require faster data sharing across parallel jobs.
- Data sharing is slow in MapReduce due to replication, serialization, and disk IO. Regarding storage system, most of the Hadoop applications, they spend more than 90% of the time doing HDFS read-write operations.



**Figure:** Interactive operations on MapReduce

# One Solution is Apache Spark

- A new general framework, which solves many of the shortcomings of MapReduce:
  - Has many other workflows, i.e. join, filter, flatMapdistinct, groupByKey, reduceByKey, sortByKey, collect, count, first...
    - (around 30 efficient distributed operations)
  - In-memory caching of data (for iterative, graph, and machine learning algorithms, etc.)
  - Original key construct: Resilient Distributed Dataset(RDD)
  - More recently added: DataFrames & DataSets
    - Different APIs for aggregate data
- It capable of leveraging the Hadoop ecosystem, e.g. HDFS, YARN, HBase, S3, ...
- Native Scala, Java, Python, and R support
- Supports interactive shells for exploratory data analysis
- Spark API is extremely simple to use
- Developed at AMPLab UC Berkeley, now by Databricks.com

# Spark


- Implements Resilient Distributed Datasets (RDDs)
- Operations on RDDs
  - Transformations: defines new dataset based on previous ones
  - Actions: starts a job to execute on cluster
- Well-designed interface to represent RDDs
  - Makes it very easy to implement transformations
  - Most Spark transformation implementation < 20 LoC

Operation	Meaning
<code>partitions()</code>	Return a list of Partition objects
<code>preferredLocations(<i>p</i>)</code>	List nodes where partition <i>p</i> can be accessed faster due to data locality
<code>dependencies()</code>	Return a list of dependencies
<code>iterator(<i>p</i>, <i>parentIters</i>)</code>	Compute the elements of partition <i>p</i> given iterators for its parent partitions
<code>partitioner()</code>	Return metadata specifying whether the RDD is hash/range partitioned

Table 3: Interface used to represent RDDs in Spark.

# Iteration in RDD

- The key idea of spark is Resilient Distributed Datasets (RDD); it supports in-memory processing computation. Data sharing in memory is 10 to 100 times faster than network and Disk.




**Figure:** Iterative operations on Spark RDD



# Interactiveness in Spark

- The key idea of spark is Resilient Distributed Datasets (RDD); it supports in-memory processing computation. Data sharing in memory is 10 to 100 times faster than network and Disk.



**Figure:** Interactive operations on Spark RDD

# Spark vs. Hadoop MapReduce

"Spark vs. Hadoop MapReduce" by Xplenty (March 11, 2019)

<https://www.xplenty.com/blog/apache-spark-vs-hadoop-mapreduce/>

- **Performance:** Spark performs better when all the data fits in the memory, especially on dedicated clusters; Hadoop MapReduce is designed for data that doesn't fit in the memory and it can run well alongside other services.
- **Ease of use:** Spark is easier to program and includes an interactive mode; Hadoop MapReduce is more difficult to program but many tools are available to make it easier.
- **Compatibility:** Spark's compatibility to data types and data sources is the same as Hadoop MapReduce.
- **Data processing:** Spark is the Swiss army knife of data processing; Hadoop MapReduce is the commando knife of batch processing.
- **Failure Tolerance :** Spark and Hadoop MapReduce both have good failure tolerance, but Hadoop MapReduce is slightly more tolerant.
- **Security :** Spark security is still in its infancy; Hadoop MapReduce has more security features and projects.
- **Maturity:** Spark maturing, Hadoop MapReduce mature

# Resilient Distributed Datasets (RDDs)

## ■ Restricted form of distributed shared memory

- Immutable, partitioned collections of records
- Allows for coarse-grained deterministic transformations (map, filter, join, ...)
- All the different processing components in Spark share the same abstraction called RDD
- As applications share the RDD abstraction, you can mix different kind of transformations to create new RDDs
- The user can control its partition
- The user can control its persistence (caching)

## ■ Efficient fault recovery using lineage

- Each RDD “remembers” how it was derived from other RDD / stable storage
- Log one operation to apply to many elements
- Recompute lost partitions on failure
- No cost if nothing fails

# RDD Actions and Transformations

- Transformations are realized when an action is called
- Transformations
  - Lazy operations applied on an RDD
  - Creates a new RDD from an existing RDD
  - Allows Spark to perform optimizations
  - e.g. map, filter, flatMap, union, intersection, distinct, reduceByKey, groupByKey
- Actions
  - Returns a value to the driver program after computation
  - e.g. reduce, collect, count, first, take, saveAsFile

# Spark Operations

<b>Transformations</b> (create a new RDD)	map filter sample <u>groupByKey</u> <u>reduceByKey</u> <u>sortByKey</u> intersection	<u>flatMap</u> union join <u>cogroup</u> cross <u>mapValues</u> <u>reduceByKey</u>
<b>Actions</b> (return results to driver program)	collect Reduce Count <u>takeSample</u> take <u>lookupKey</u>	first take <u>takeOrdered</u> <u>countByKey</u> save <u>foreach</u>

# Sample Spark transformations

- `map(func)`: Return a new distributed dataset formed by passing each element of the source through a function `func`.
- `filter(func)`: Return a new dataset formed by selecting those elements of the source on which `func` returns true
- `union(otherDataset)`: Return a new dataset that contains the union of the elements in the source dataset and the argument.
- `intersection(otherDataset)`: Return a new RDD that contains the intersection of elements in the source dataset and the argument.
- `distinct([numTasks])`: Return a new dataset that contains the distinct elements of the source dataset
- `join(otherDataset, [numTasks])`: When called on datasets of type  $(K, V)$  and  $(K, W)$ , returns a dataset of  $(K, (V, W))$  pairs with all pairs of elements for each key. Outer joins are supported through `leftOuterJoin`, `rightOuterJoin`, and `fullOuterJoin`.

Source: <https://spark.apache.org/docs/latest/programming-guide.html>



# Sample Spark Actions

- `reduce(func)`: Aggregate the elements of the dataset using a function `func` (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
- `collect()`: Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
- `count()`: Return the number of elements in the dataset.

Remember: Actions cause calculations to be performed; transformations just set things up (lazy evaluation)

Source: <https://spark.apache.org/docs/latest/programming-guide.html>



# Spark 2.0

---

- Note that, before Spark 2.0, the main programming interface of Spark was the Resilient Distributed Dataset (RDD).
- After Spark 2.0, RDDs are replaced by **Dataset**, which is strongly-typed like an RDD, but with richer optimizations under the hood. The RDD interface is still supported, and you can get a more detailed reference at the RDD programming guide.

# Spark Ecosystem

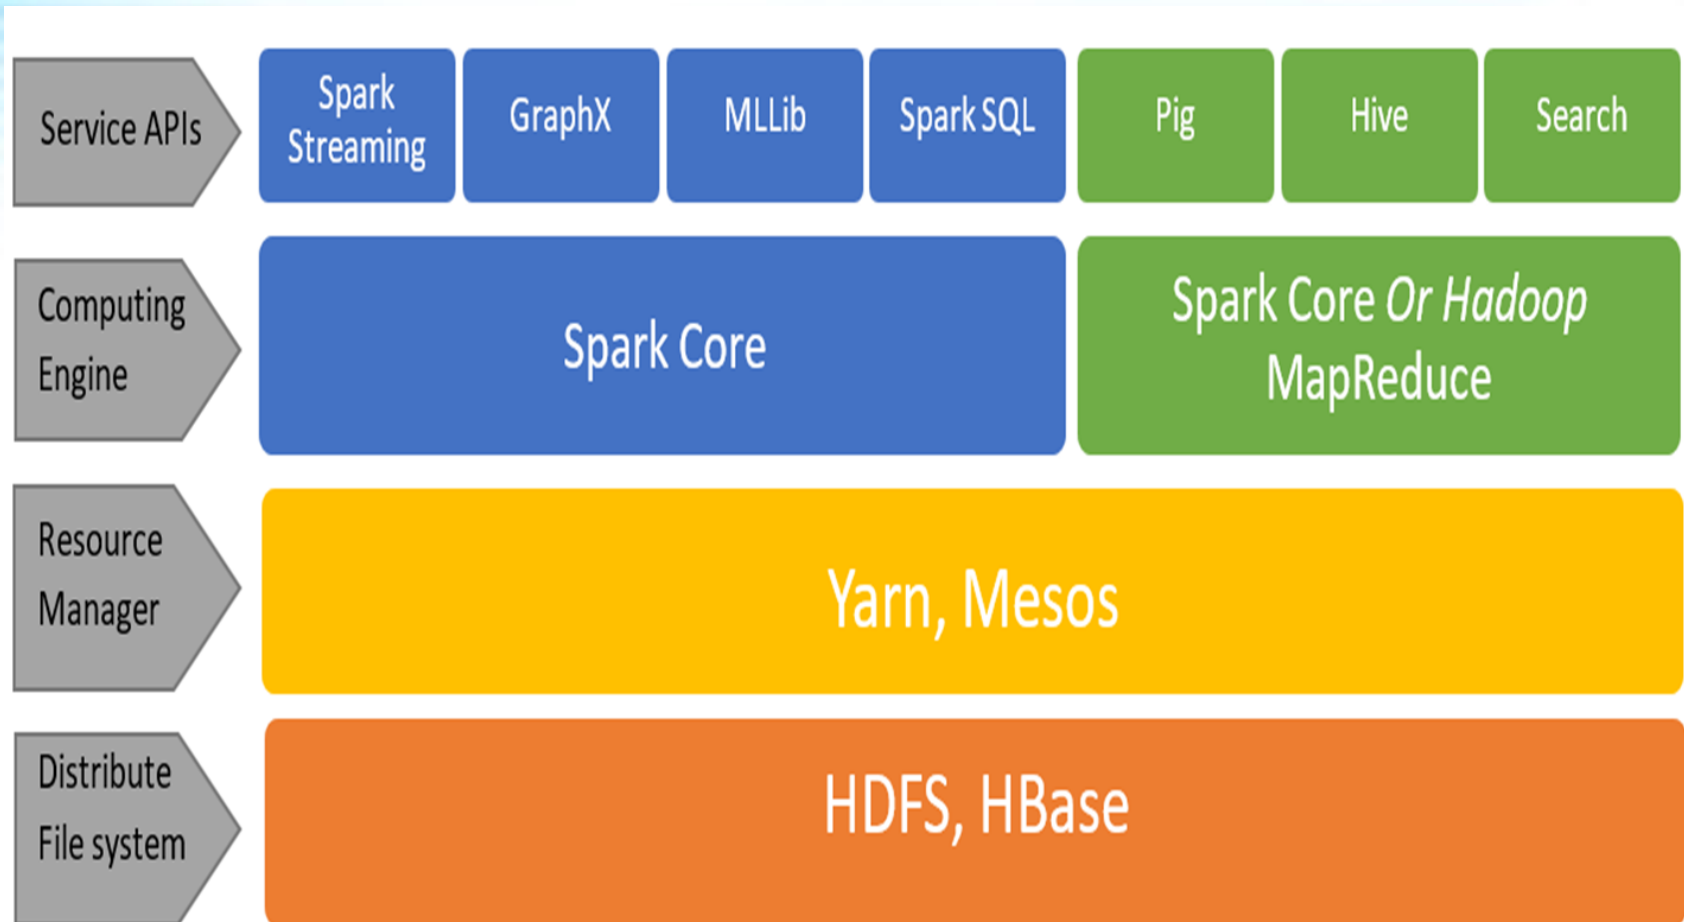


Figure 1 - Spark Context

# Architecture of Spark Applications

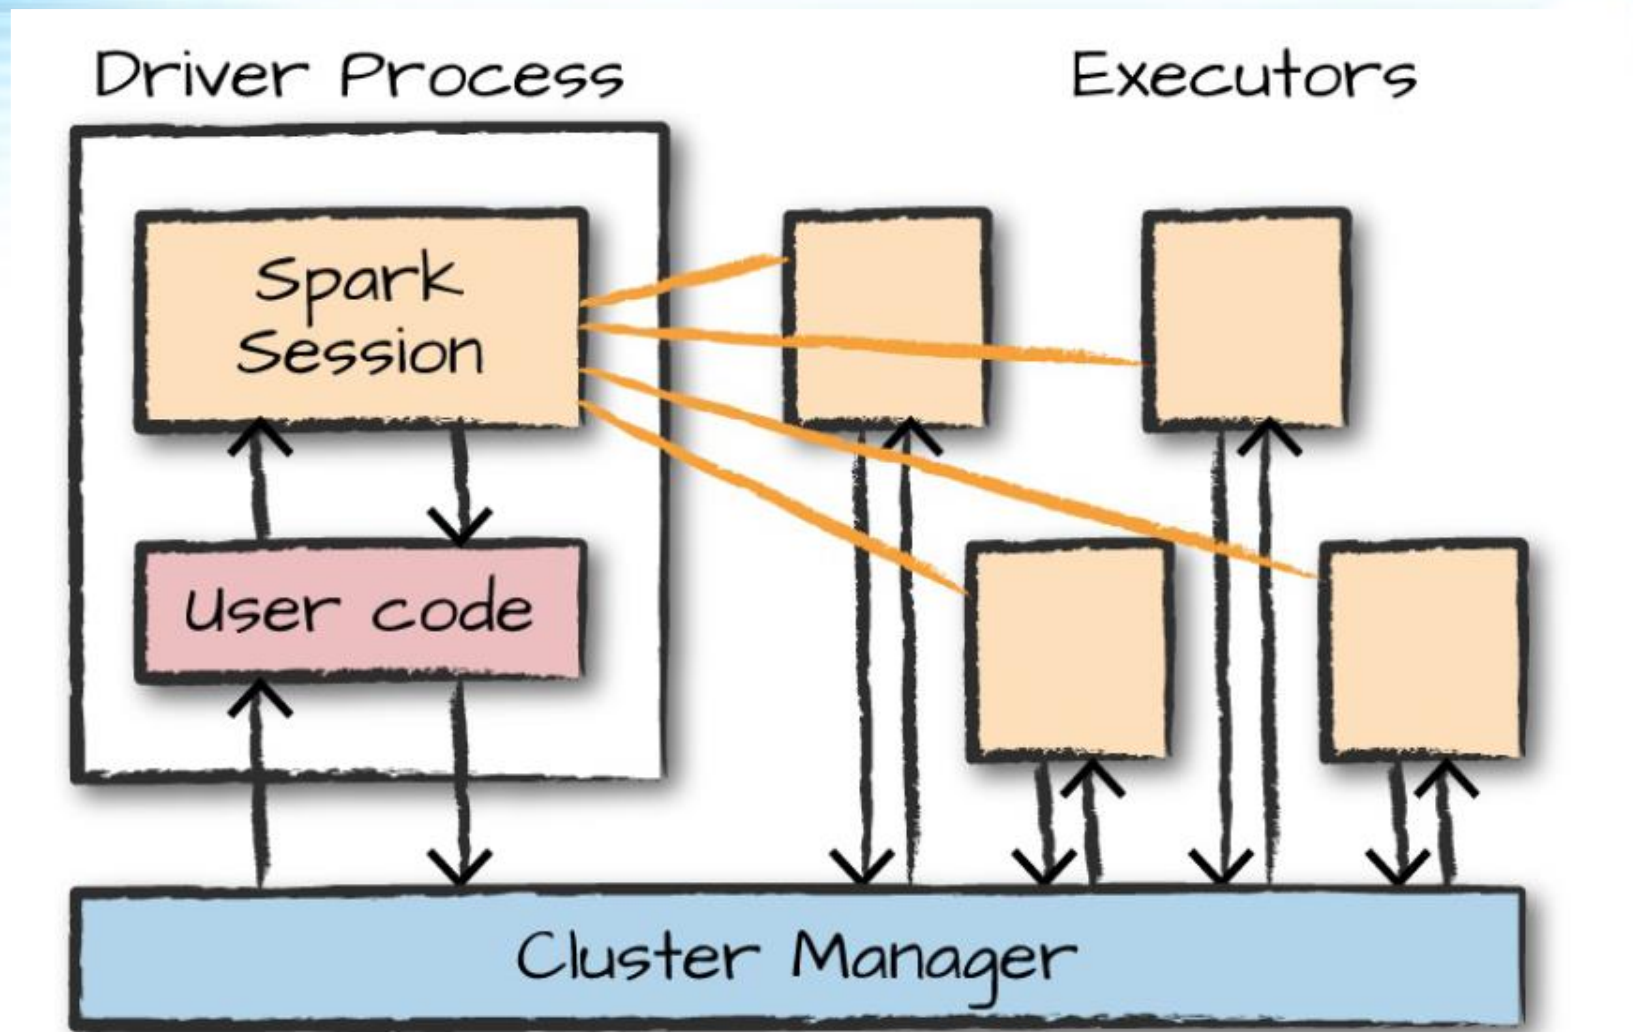
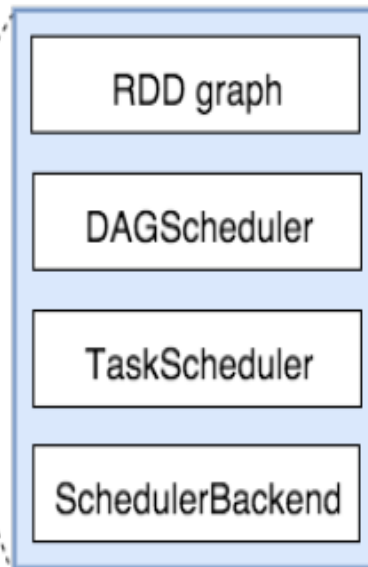


Figure 2-1. The architecture of a Spark Application


User Program

```
val sc = new SparkContext(conf) {  
  val rdd = sc.cassandraTable(...)  
    .map(...)  
    .filter(...)  
    .keyBy(...)  
    .reduceByKey(...)  
    .cache()  
}
```

Driver



Executor



# RDD, DataFrame, Dataset

- Data Import and low-level coding
- Application programming

RDDs

- Table-based functionality
- SQL-style query access  
*(R or Python users will default to this as well)*

DataFrames



- More intensive application programming in Java
- Require Class declaration and definition

Datasets

# DataFrame & Dataset

<https://data-flair.training/blogs/apache-spark-rdd-vs-dataframe-vs-dataset/>

<https://databricks.com/blog/2016/07/14/a-tale-of-three-apache-spark-apis-rdds-dataframes-and-datasets.html>

## ■ DataFrame:

- Unlike an RDD, data organized into named columns, e.g. a table in a relational database.
- Imposes a structure onto a distributed collection of data, allowing higher-level abstraction

## ■ Dataset:

- Extension of DataFrame API which provides type-safe, object-oriented programming interface (compile-time error detection)

Both built on Spark SQL engine & use Catalyst to generate optimized logical and physical query plan; both can be converted to an RDD



# Dataframe Transformation





Figure 2-10. The entire DataFrame transformation flow



# RDDs vs. DataFrames

- RDDs provide a low level interface into Spark
- DataFrames have a schema
- DataFrames are cached and optimized by Spark
- DataFrames are built on top of the RDDs and the core Spark API



# API distinction: typing

- Python & R don't have compile-time type safety checks, so *only* support DataFrame
  - Error detection only at runtime
- Java & Scala support compile-time type safety checks, so support *both* DataSet and DataFrame
  - Dataset APIs are all expressed as lambda functions and JVM typed objects
  - any mismatch of typed-parameters will be detected at compile time.

<https://databricks.com/blog/2016/07/14/a-tale-of-three-apache-spark-apis-rdds-dataframes-and-datasets.html>

# Spark in the Real World (I)

- Uber – the online taxi company gathers terabytes of event data from its mobile users every day.
  - By using Kafka, Spark Streaming, and HDFS, to build a continuous ETL pipeline
  - Convert raw unstructured event data into structured data as it is collected
  - Uses it further for more complex analytics and optimization of operations
  
- Pinterest – Uses a Spark ETL pipeline
  - Leverages Spark Streaming to gain immediate insight into how users all over the world are engaging with Pins—in real time.
  - Can make more relevant recommendations as people navigate the site
  - Recommends related Pins
  - Determine which products to buy, or destinations to visit

# Spark in the Real World (II)

---

Here are Few other Real World Use Cases:

- **Conviva – 4 million video feeds per month**
  - This streaming video company is second only to YouTube.
  - Uses Spark to reduce customer churn by optimizing video streams and managing live video traffic
  - Maintains a consistently smooth, high quality viewing experience.
- **Capital One – is using Spark and data science algorithms to understand customers in a better way.**
  - Developing next generation of financial products and services
  - Find attributes and patterns of increased probability for fraud
- **Netflix – leveraging Spark for insights of user viewing habits and then recommends movies to them.**
  - User data is also used for content creation

# Spark: when not to use

- Even though Spark is versatile, that doesn't mean Spark's in-memory capabilities are the best fit for all use cases:
  - For many simple use cases Apache MapReduce and Hive might be a more appropriate choice
  - Spark was not designed as a multi-user environment
  - Spark users are required to know that memory they have is sufficient for a dataset
  - Adding more users adds complications, since the users will have to coordinate memory usage to run code

# Apache Spark Extensions “on top” of core

---

- Spark SQL
- Spark Streaming – stream processing of live datastreams
- GraphX – graph manipulation
  - extends Spark RDD with Graph abstraction: a directed multigraph with properties attached to each vertex and edge.
- MLlib – machine learning